
Real-time Systems

Chapter 9:

Resource Access Protocol

Ngo Lam Trung
Dept. of Computer Engineering

Contents

- ❑ Introduction
- ❑ The priority inversion phenomenon
- ❑ Solutions for priority inversion

Resource constraint

❑ Resource

- ❑ Any software structure that can be used by the process to advance its execution
- ❑ Ex: data structure, variables, main memory area, a file, a piece of program, a set of registers of a peripheral device

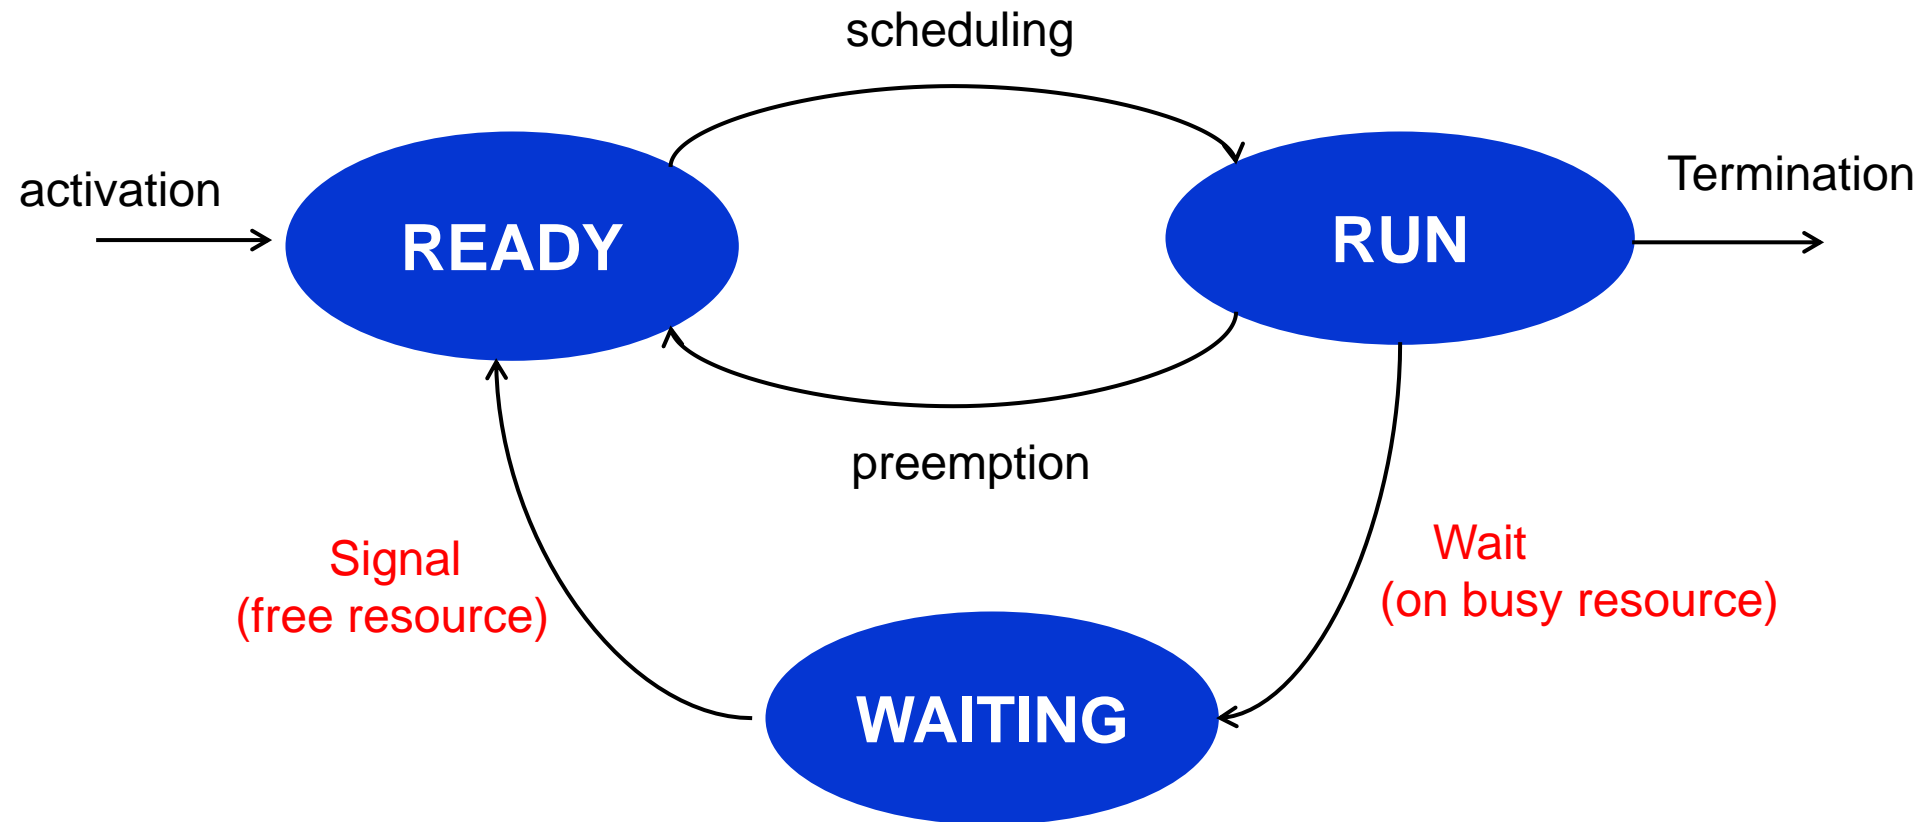
❑ Many shared resources do not allow simultaneous access

→ require **mutual exclusion**

❑ Critical section


- ❑ A piece of code under mutual exclusion constraints
- ❑ Tasks entering critical section have to wait until no other task is holding the resource

Waiting state caused by resource constraint



Example of blocking on exclusive resource


❑ Scheduling with preemption



Problem: the task with higher priority has to wait for the task with lower priority

The priority inversion phenomenon

- ❑ J3 enters critical section first
- ❑ J1 is blocked, has to wait until J3 signal the resource
- ❑ J2 preempts J3 → J1 has to wait for J2



Problems

- ❑ The task with higher priority has to wait for the task with lower priority
- ❑ Blocking time is unbounded → the system is not predictable.
- ❑ Example of priority inversion: Mars Pathfinder 1997
 - ❑ CPU: RAD6000 20MHz (\$200K-\$300K)
 - ❑ OS: VxWork
 - ❑ Experienced CPU reset upon touching down on Mars, debugging on Earth detected priority inversion, fixed by new firmware upload.

Problems

❑ Solutions

- ❑ Non-preemptive Protocol
- ❑ Highest Locker Priority Protocol
- ❑ Priority Inheritance Protocol
- ❑ Priority Ceiling Protocol
- ❑ Stack Resource Policy

Terminology & assumptions(1)

- ❑ Periodic task set $\Gamma = \{\tau_1, \tau_2, \dots, \tau_n\}$
 - ❑ $\tau_i = (C_i, T_i)$
 - ❑ Relative deadline $D_i = T_i$
- ❑ Resources R_1, \dots, R_m
 - ❑ Each R_k is guarded by semaphore S_k
- ❑ J_i : a job of τ_i
- ❑ P_i : nominal priority of τ_i
- ❑ $p_i \geq P_i$: active priority of τ_i (initially set to P_i)
- ❑ $z_{i,j}$: j -th critical section of J_i
- ❑ $d_{i,j}$: duration of $z_{i,j}$
- ❑ $S_{i,j}$: the semaphore guarding $z_{i,j}$
- ❑ $R_{i,j}$: the resource used in $z_{i,j}$
- ❑ Notation $z_{i,j} \subset z_{i,k}$ means $z_{i,j}$ is entirely contained in $z_{i,k}$.

Terminology & assumptions (2)

□ Assumptions

- J_1, \dots, J_n are listed in decreasing order of P_i
- Jobs don't suspend themselves.
- The critical sections used by any task are properly nested.

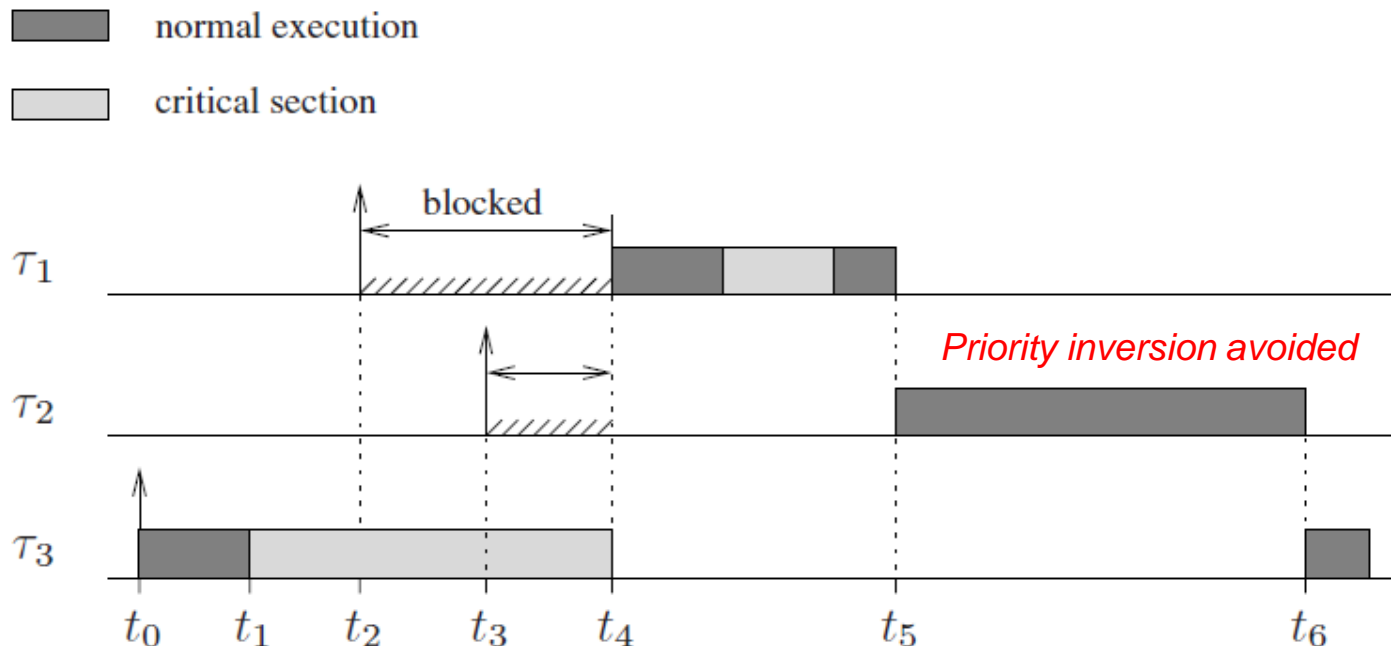
$$Z_{i,j} \subset Z_{i,k} \text{ or } Z_{i,k} \subset Z_{i,j} \text{ or } Z_{i,j} \cap Z_{i,k} = \emptyset$$

- Critical sections are guarded by binary semaphores.

The simplest: Non-preemptive Protocol


- ❑ Block all other tasks whenever a task enters a critical section
- ❑ The dynamic priority of the running task is raised to the highest level

$$p_i(R_k) = \max_h \{P_h\}$$



The simplest: Non-preemptive Protocol (NPP)

- ❑ Pros: simple
- ❑ Cons: unnecessary blocking



Blocking time of Non-preemptive Protocol

- ❑ Given the task T_i , the set of critical sections that can block T_i

$$\gamma_i = \{Z_{j,k} \mid P_j < P_i, k = 1, \dots, m\}$$

- ❑ The maximum blocking time is

$$B_i = \max\{d_{j,k} - 1 \mid Z_{j,k} \in \gamma_i\}.$$

➔ Duration of the longest critical section that can block T_i

Highest Locker Priority Protocol (HLP)

- ❑ Improves NPP: raising the priority of the task entering a critical section to the highest priority among the tasks sharing that resource.
- ❑ When a task enters resource R_k , its dynamic priority is raised to


$$p_i(R_k) = \max_h \{P_h \mid \tau_h \text{ uses } R_k\}$$

- ❑ When the task exits the resource, its dynamic priority is reset to the nominal value P_i
- ❑ Priority ceiling can be computed offline

$$C(R_k) \stackrel{\text{def}}{=} \max_h \{P_h \mid \tau_h \text{ uses } R_k\}$$

Highest Locker Priority Protocol

Example



HLP Blocking time

- The set of critical instants that can block task T_i

$$\gamma_i = \{Z_{j,k} \mid (P_j < P_i) \text{ and } C(R_k) \geq P_i\}$$

- Hence, maximum blocking time is

$$B_i = \max_{j,k} \{\delta_{j,k} - 1 \mid Z_{j,k} \in \gamma_i\}$$

- Problem: what if critical section is access in only one branch of a conditional statement?


Priority Inheritance Protocol

- ❑ Modify the priority of tasks in critical sections
- ❑ When a task blocks higher-priority tasks, it temporarily ***inherits*** the highest priority of the blocked tasks.
 - ❑ Prevents preemption of medium-priority tasks

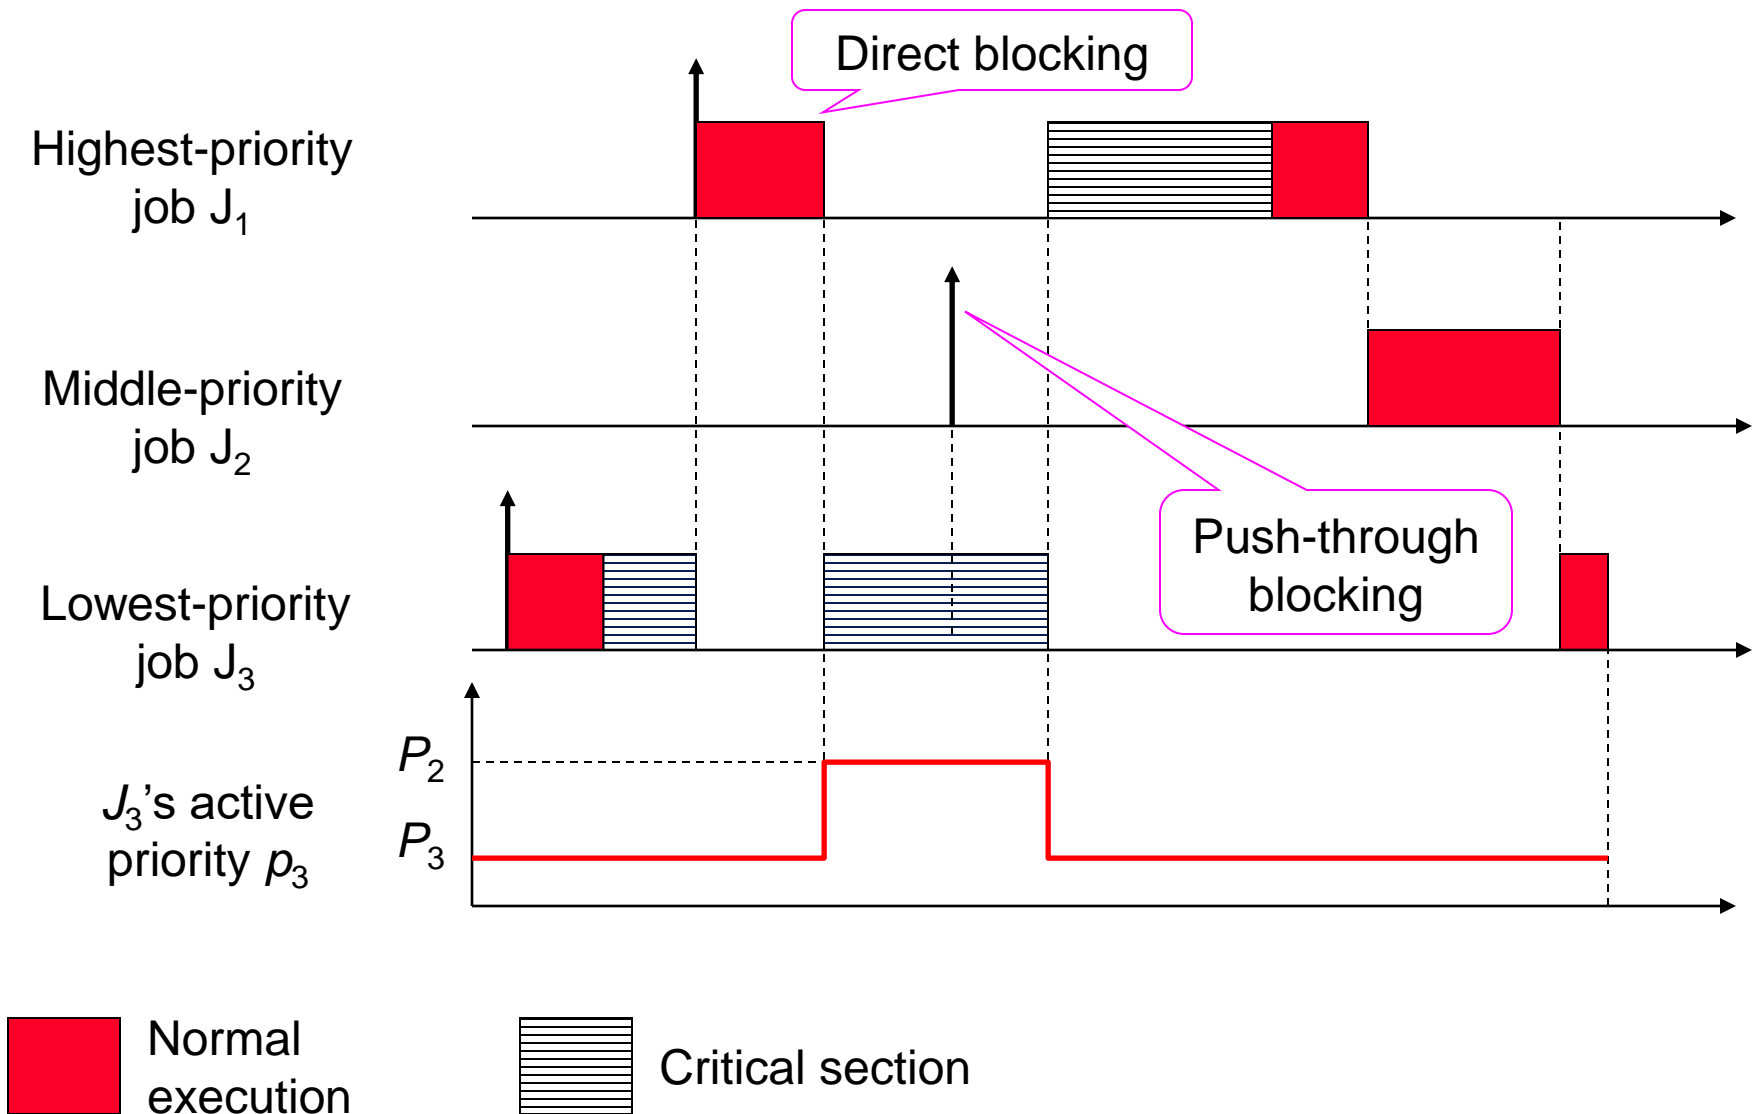
Protocol definition

- ❑ Jobs are scheduled based on their active priorities
- ❑ When the higher-priority job J_{high} is blocked on a semaphore because the lower-priority job J_{low} is in execution of its critical section, **the active priority p_{high} of J_{high} is inherited** to that of J_{low} .
- ❑ The rest of the critical section of J_{low} is executed with the active priority p_{high} .
- ❑ In case the medium-priority job J_{medium} activates, it cannot preempt the execution of J_{low} → Unbounded priority inversion is avoided.
- ❑ Priority inheritance is transitive; if a job J_3 blocks a job J_2 , and J_2 blocks a job J_1 , then J_3 inherits the priority of J_1 via J_2 .

Example




Direct blocking & Push-through blocking



PIP with nested critical sections

- ❑ When the blocking job J_k exits the critical section, the blocked job with the highest priority is awakened.
- ❑ J_k replaces its active priority p_k by nominal priority P_k if no other jobs are blocked by J_k , or by the highest priority of the tasks blocked by J_k



Transitive
priority
inheritance


Properties

- ❑ Push-through blocking to job J_i occurs only if the semaphore is accessed by a job J_{low} with $p_{\text{low}} < p_i$ and by a job J_{high} with p_{high} that can be equal or higher than p_i
- ❑ Transitive priority inheritance can occur only in the presence of nested critical sections.
- ❑ If there are n lower-priority jobs that can block a job J_i , then J_i can be blocked at most the duration of n critical sections.
- ❑ If there are m distinct semaphores that can block a job J_i , then J_i can be blocked for at most the duration of m critical sections.

Properties

- Under the priority inheritance protocol, a job J can be blocked for at most the duration of $\min(n, m)$ critical sections.
 - n is the number of lower-priority jobs that could block J
 - m is the number of distinct semaphores that can be used to block J
- ➔ *The maximum blocking time for any task J is bounded*



Remaining problem 1: Chained blocking



→ J_1 can be blocked several times

Remaining problem (2): Deadlock


■ normal execution
■ critical section



➔ Deadlock caused as J_2 enters the nested critical session

Priority Ceiling Protocol

- ❑ Extends the Priority Inheritance Protocol
- ❑ Assign each semaphore a ceiling priority, equal to the priority of the highest-priority task that can lock it.
- ❑ Provided a critical section contains several semaphores, a job J can enter the critical section only when its priority is higher than all priority ceilings of the semaphores already locked by other jobs.



Protocol definition (1)

- ❑ S_k : an arbitrary semaphore
- ❑ $C(S_k)$: priority ceiling of S_k

$$C(S_k) \stackrel{\text{def}}{=} \max_i \{P_i \mid S_k \in \sigma_i\}.$$

This value can be computed offline

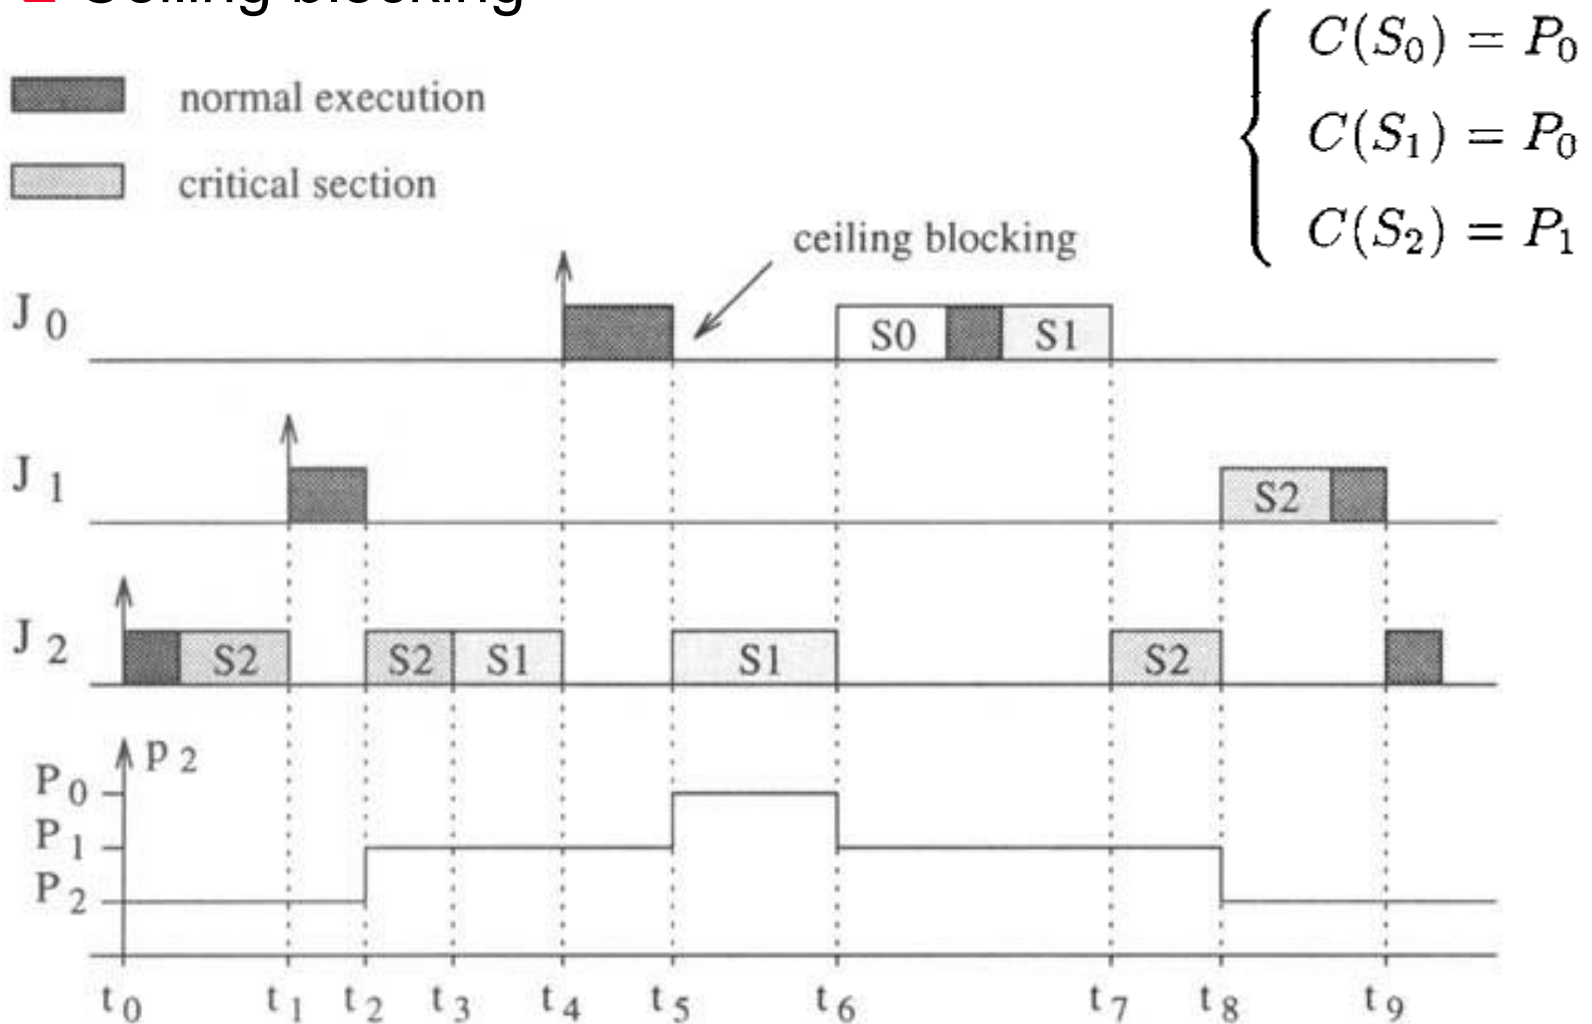
- ❑ J_i : the job with the highest priority in ready queue
- ❑ P_i : the priority of J_i
- ❑ S^* : semaphore with the highest priority ceiling among all the semaphores currently locked by jobs other than J_i

Protocol definition (2)

- ❑ When J_i is about to enter a critical section guarded by semaphore S_k ,
 - ❑ If $P_i \leq C(S^*)$
 - locking on S_k is denied, &
 - J_i is blocked on semaphore S^* by the job holding the lock on S^* .
 - ❑ If $P_i > C(S^*)$
 - J_i locks on S_k and continue execution
- ❑ When J_i is blocked on a semaphore S ,
 - ❑ The job J_k locking on S inherits the priority p_i
 - ❑ Generally, a task inherits the highest priority of the jobs blocked by it.
- ❑ When J_k exits a critical section & unlocks the semaphore,
 - ❑ If there are blocked jobs, then p_k is the highest active priority of the jobs blocked by J_k
 - ❑ Otherwise, p_k is restored to P_k


Example

❑ Ceiling blocking



Ceiling blocking

- ❑ A task is blocked by the protocol because of the priority ceiling condition
- ❑ Necessary to avoid chained blocking and deadlock



This will never happen with PCP

Properties of the protocol (2)

- ❑ The Priority Ceiling Protocol prevents deadlocks.
- ❑ Under the Priority Ceiling Protocol, a job J_i can be blocked for at most the duration of one critical section.

- ➔ Reduce blocking time
- ➔ Avoid unnecessary high-priority tasks blocking
- ➔ Avoid deadlock

Comparison

	priority	Num. of blocking	pessimism	blocking instant	transparency	deadlock prevention	implementation
NPP	any	1	high	on arrival	YES	YES	easy
HLP	fixed	1	medium	on arrival	NO	YES	easy
PIP	fixed	α_i	low	on access	YES	NO	hard
PCP	fixed	1	medium	on access	NO	YES	medium
SRP	any	1	medium	on arrival	NO	YES	easy

SRP (Stack Resource Protocol): for student's further reading