
Chapter 3: Arithmetic for Computers

Ngo Lam Trung

[with materials from *Computer Organization and Design, 4th Edition*,
Patterson & Hennessy, © 2008, MK
and M.J. Irwin's presentation, PSU 2008]

Content

- ❑ Integer arithmetic
- ❑ Floating point number representation and arithmetic

Addition and subtraction

□ Addition

- Similar to what you do to add two numbers manually
- Digits are added bit by bit from right to left
- Carries passed to the next digit to the left

□ Subtraction

- Negate the second operand then add to the first operand

$$\begin{array}{r} + \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0111_{\text{two}} = 7_{\text{ten}} \\ \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0110_{\text{two}} = 6_{\text{ten}} \\ \hline \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 1101_{\text{two}} = 13_{\text{ten}} \end{array}$$

Examples

- ❑ All numbers are 8-bit signed integer

$$12 + 8 =$$

$$122 + 8 =$$

$$122 + 80 =$$

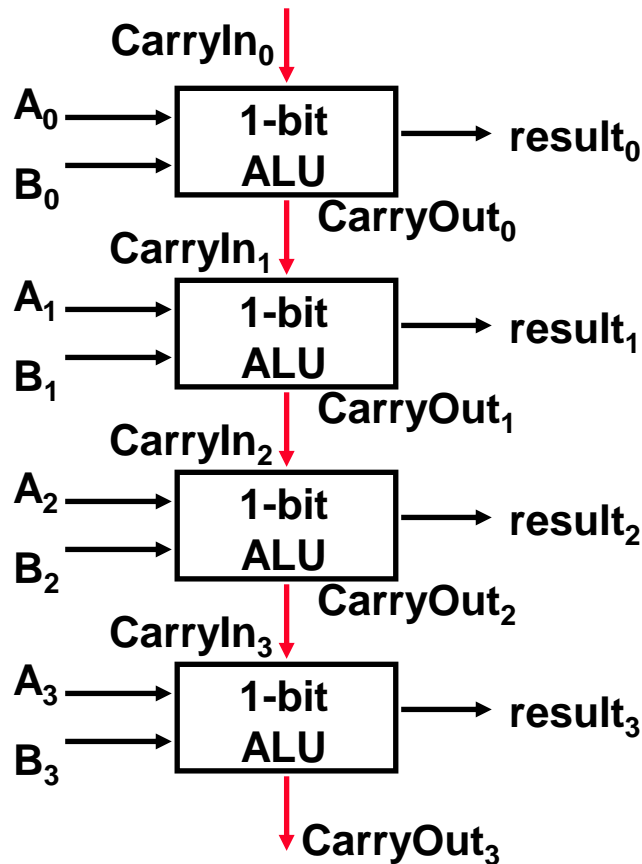
Dealing with Overflow

- ❑ Overflow occurs when the result of an operation cannot be represented in 32-bits, i.e., when the sign bit contains a **value** bit of the result and not the proper **sign** bit
- ❑ **When adding operands with different signs or when subtracting operands with the same sign, overflow can *never* occur**

Operation	Operand A	Operand B	Result indicating overflow
A + B	≥ 0	≥ 0	< 0
A + B	< 0	< 0	≥ 0
A - B	≥ 0	< 0	< 0
A - B	< 0	≥ 0	≥ 0

Adder implementation

❑ N-bit ripple-carry adder



Performance depends
on data length

➔ Performance is low

Making addition faster: infinite hardware

❑ Parallelize the adder with the cost of hardware

❑ Given the addition:

$$a_{n-1}a_{n-2} \dots a_1a_0 + b_{n-1}b_{n-2} \dots b_1b_0$$

❑ Let c_i is the carry at bit i

$$c_2 = (b_1 \cdot c_1) + (a_1 \cdot c_1) + (a_1 \cdot b_1)$$

$$c_1 = (b_0 \cdot c_0) + (a_0 \cdot c_0) + (a_0 \cdot b_0)$$

Find c_2 from a_0, b_0, a_1, b_1 ?

Making addition faster: Carry Look-ahead

❑ Approach

- ❑ Make hardwired 4 bit adder → fast and simple enough
- ❑ Develop a carry look-ahead unit to calculate the carry bit before finishing the addition

❑ At bit i

$$\begin{aligned} c_{i+1} &= (b_i \cdot c_i) + (a_i \cdot c_i) + (a_i \cdot b_i) \\ &= (a_i \cdot b_i) + (a_i + b_i) \cdot c_i \end{aligned}$$

❑ Denote

$$\begin{aligned} g_i &= a_i \cdot b_i \\ p_i &= a_i + b_i \end{aligned}$$

❑ Then

$$c_{i+1} = g_i + p_i \cdot c_i$$

Carry look-ahead

❑ With 4-bit adder

$$c1 = g0 + (p0 \cdot c0)$$

$$c2 = g1 + (p1 \cdot g0) + (p1 \cdot p0 \cdot c0)$$


$$c3 = g2 + (p2 \cdot g1) + (p2 \cdot p1 \cdot g0) + (p2 \cdot p1 \cdot p0 \cdot c0)$$

$$c4 = g3 + (p3 \cdot g2) + (p3 \cdot p2 \cdot g1) + (p3 \cdot p2 \cdot p1 \cdot g0) \\ + (p3 \cdot p2 \cdot p1 \cdot p0 \cdot c0)$$

- ➔ All carry bits can be calculated after 3 gate delay
- ➔ All result bits can be calculated after maximum of 4 gate delay
- ➔ How to implement bigger adder?

Carry look-ahead

- For 16-bit adder → fast C1, C2, C3, C4 is needed



Carry look-ahead

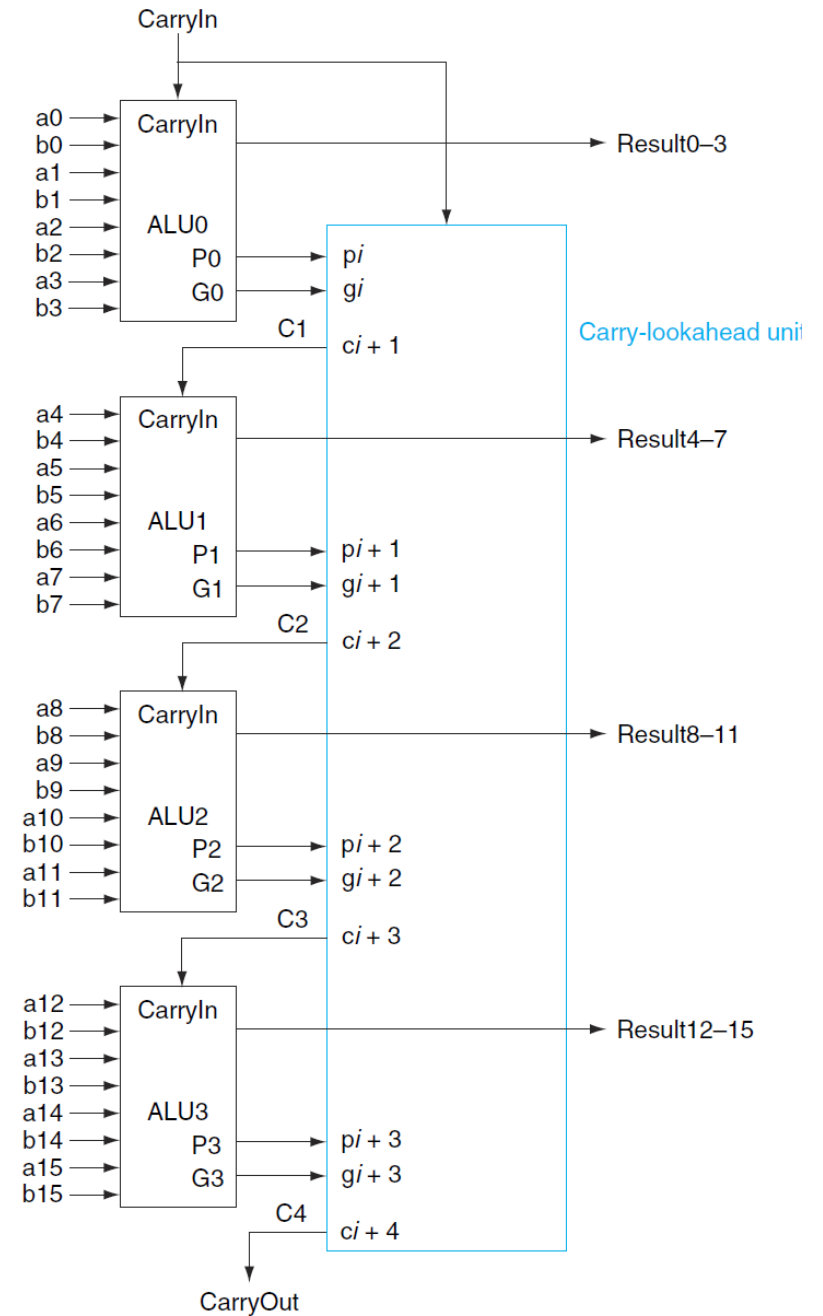
□ Denote

$$\begin{aligned}P_0 &= p_3 \cdot p_2 \cdot p_1 \cdot p_0 & G_0 &= g_3 + (p_3 \cdot g_2) + (p_3 \cdot p_2 \cdot g_1) + (p_3 \cdot p_2 \cdot p_1 \cdot g_0) \\P_1 &= p_7 \cdot p_6 \cdot p_5 \cdot p_4 & G_1 &= g_7 + (p_7 \cdot g_6) + (p_7 \cdot p_6 \cdot g_5) + (p_7 \cdot p_6 \cdot p_5 \cdot g_4) \\P_2 &= p_{11} \cdot p_{10} \cdot p_9 \cdot p_8 & G_2 &= g_{11} + (p_{11} \cdot g_{10}) + (p_{11} \cdot p_{10} \cdot g_9) + (p_{11} \cdot p_{10} \cdot p_9 \cdot g_8) \\P_3 &= p_{15} \cdot p_{14} \cdot p_{13} \cdot p_{12} & G_3 &= g_{15} + (p_{15} \cdot g_{14}) + (p_{15} \cdot p_{14} \cdot g_{13}) + (p_{15} \cdot p_{14} \cdot p_{13} \cdot g_{12})\end{aligned}$$

□ Then big-carry bits can be calculated fast

$$\begin{aligned}C_1 &= G_0 + (P_0 \cdot c_0) \\C_2 &= G_1 + (P_1 \cdot G_0) + (P_1 \cdot P_0 \cdot c_0) \\C_3 &= G_2 + (P_2 \cdot G_1) + (P_2 \cdot P_1 \cdot G_0) + (P_2 \cdot P_1 \cdot P_0 \cdot c_0) \\C_4 &= G_3 + (P_3 \cdot G_2) + (P_3 \cdot P_2 \cdot G_1) + (P_3 \cdot P_2 \cdot P_1 \cdot G_0) \\&\quad + (P_3 \cdot P_2 \cdot P_1 \cdot P_0 \cdot c_0)\end{aligned}$$

16-bit Adder



Exercise

- ❑ Determine g_i, p_i, G_i, P_i when adding the two 16-bit numbers

$$a = 0001\ 1010\ 0011\ 0011$$

$$b = 1110\ 0101\ 1110\ 1011$$

- ❑ Calculate c_{15}

Exercise

□ p_i, g_i

$$g_i = a_i \cdot b_i$$
$$p_i = a_i + b_i$$

a:	0001	1010	0011	0011
b:	1110	0101	1110	1011
g_i :	0000	0000	0010	0011
p_i :	1111	1111	1111	1011

$$P_3 = 1 \cdot 1 \cdot 1 \cdot 1 = 1$$

$$P_2 = 1 \cdot 1 \cdot 1 \cdot 1 = 1$$

$$P_1 = 1 \cdot 1 \cdot 1 \cdot 1 = 1$$

$$P_0 = 1 \cdot 0 \cdot 1 \cdot 1 = 0$$

$$G_0 = g_3 + (p_3 \cdot g_2) + (p_3 \cdot p_2 \cdot g_1) + (p_3 \cdot p_2 \cdot p_1 \cdot g_0)$$
$$= 0 + (1 \cdot 0) + (1 \cdot 0 \cdot 1) + (1 \cdot 0 \cdot 1 \cdot 1) = 0 + 0 + 0 + 0 = 0$$

$$G_1 = g_7 + (p_7 \cdot g_6) + (p_7 \cdot p_6 \cdot g_5) + (p_7 \cdot p_6 \cdot p_5 \cdot g_4)$$
$$= 0 + (1 \cdot 0) + (1 \cdot 1 \cdot 1) + (1 \cdot 1 \cdot 1 \cdot 0) = 0 + 0 + 1 + 0 = 1$$

$$G_2 = g_{11} + (p_{11} \cdot g_{10}) + (p_{11} \cdot p_{10} \cdot g_9) + (p_{11} \cdot p_{10} \cdot p_9 \cdot g_8)$$
$$= 0 + (1 \cdot 0) + (1 \cdot 1 \cdot 0) + (1 \cdot 1 \cdot 1 \cdot 0) = 0 + 0 + 0 + 0 = 0$$

$$G_3 = g_{15} + (p_{15} \cdot g_{14}) + (p_{15} \cdot p_{14} \cdot g_{13}) + (p_{15} \cdot p_{14} \cdot p_{13} \cdot g_{12})$$
$$= 0 + (1 \cdot 0) + (1 \cdot 1 \cdot 0) + (1 \cdot 1 \cdot 1 \cdot 0) = 0 + 0 + 0 + 0 = 0$$

Exercise

□ c_{15} is actually C_4

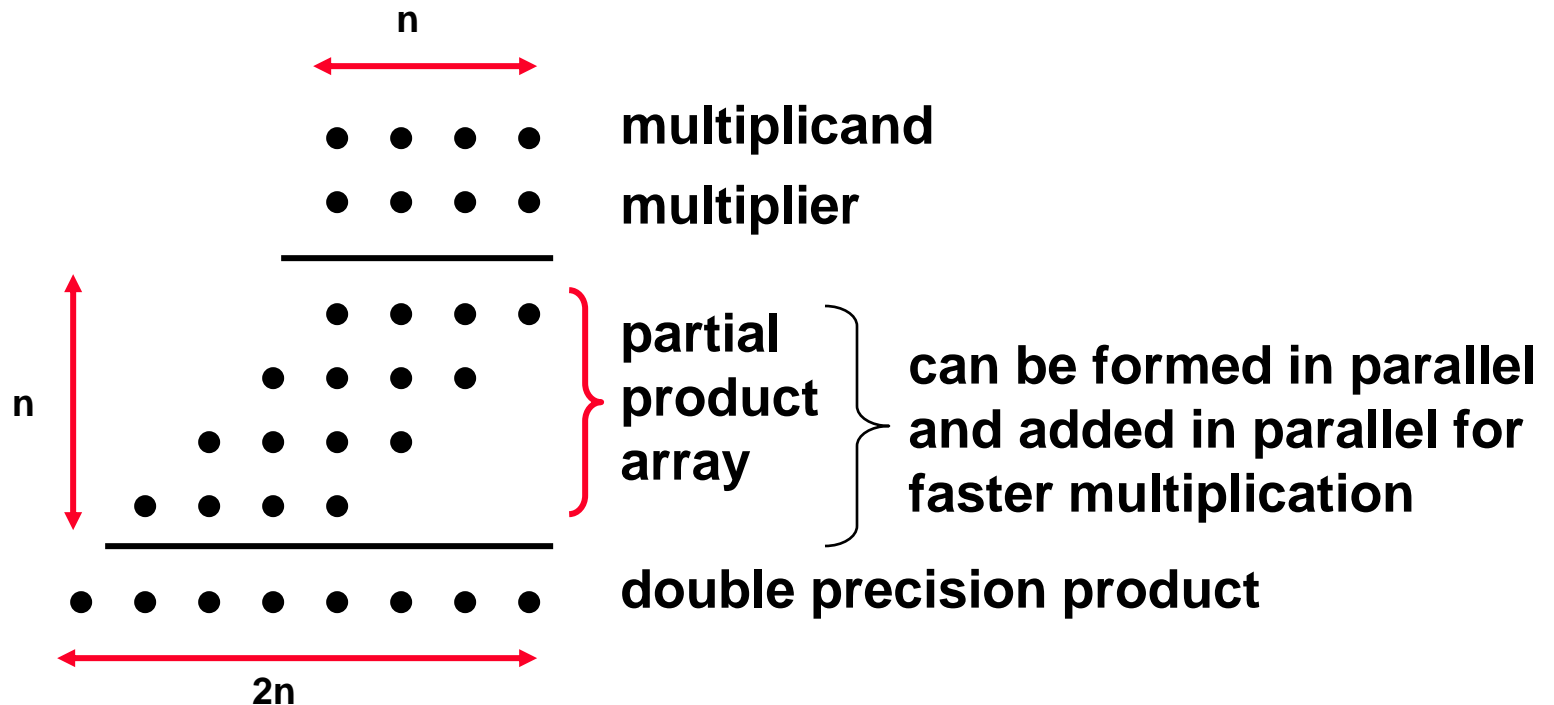
$$\begin{aligned} C_4 &= G_3 + (P_3 \cdot G_2) + (P_3 \cdot P_2 \cdot G_1) + (P_3 \cdot P_2 \cdot P_1 \cdot G_0) \\ &\quad + (P_3 \cdot P_2 \cdot P_1 \cdot P_0 \cdot c_0) \\ &= 0 + (1 \cdot 0) + (1 \cdot 1 \cdot 1) + (1 \cdot 1 \cdot 1 \cdot 0) + (1 \cdot 1 \cdot 1 \cdot 0 \cdot 0) \\ &= 0 + 0 + 1 + 0 + 0 = 1 \end{aligned}$$

Exercise

- ❑ Compare performance of 16-bit ripple carry and 16-bit carry look-ahead adders, assuming delay of all logic gates are equal?

Multiply

- ❑ Binary multiplication is just a *bunch* of right shifts and adds

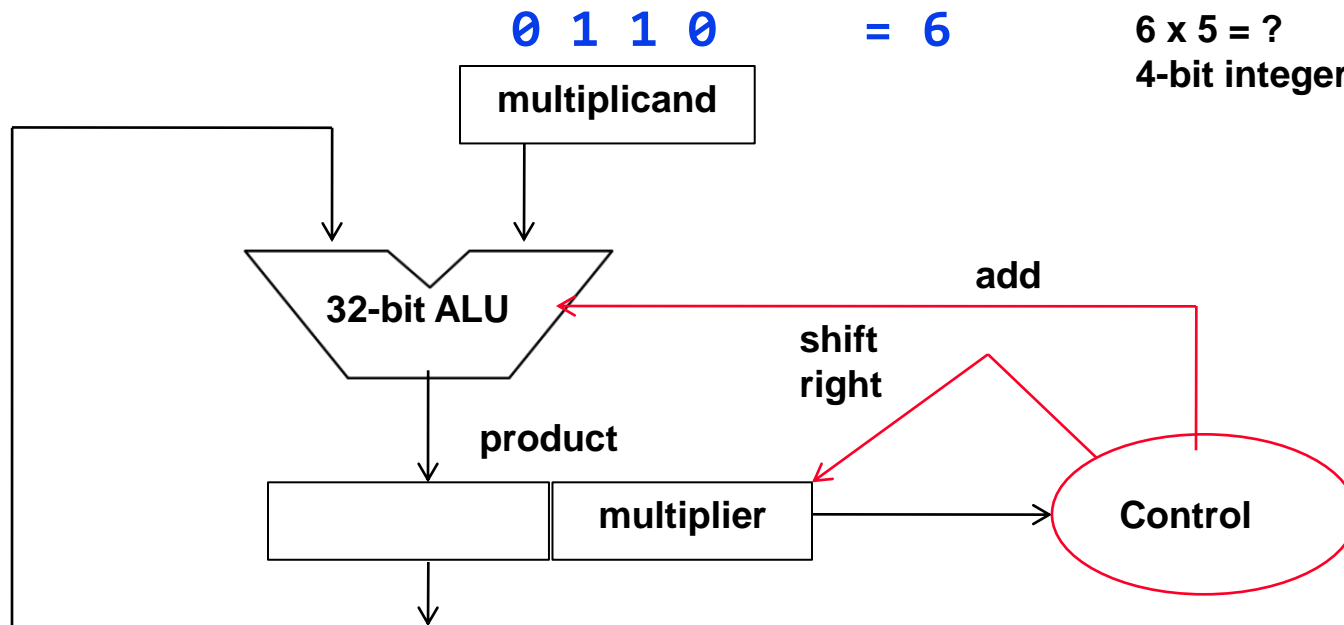


n-bit multiplicand and multiplier → 2n-bit product

Example

Multiplicand		1000 _{ten}
Multiplier	x	1001 _{ten}
		<hr/>
		1000
		0000
		0000
		1000
		<hr/>
Product		1001000 _{ten}

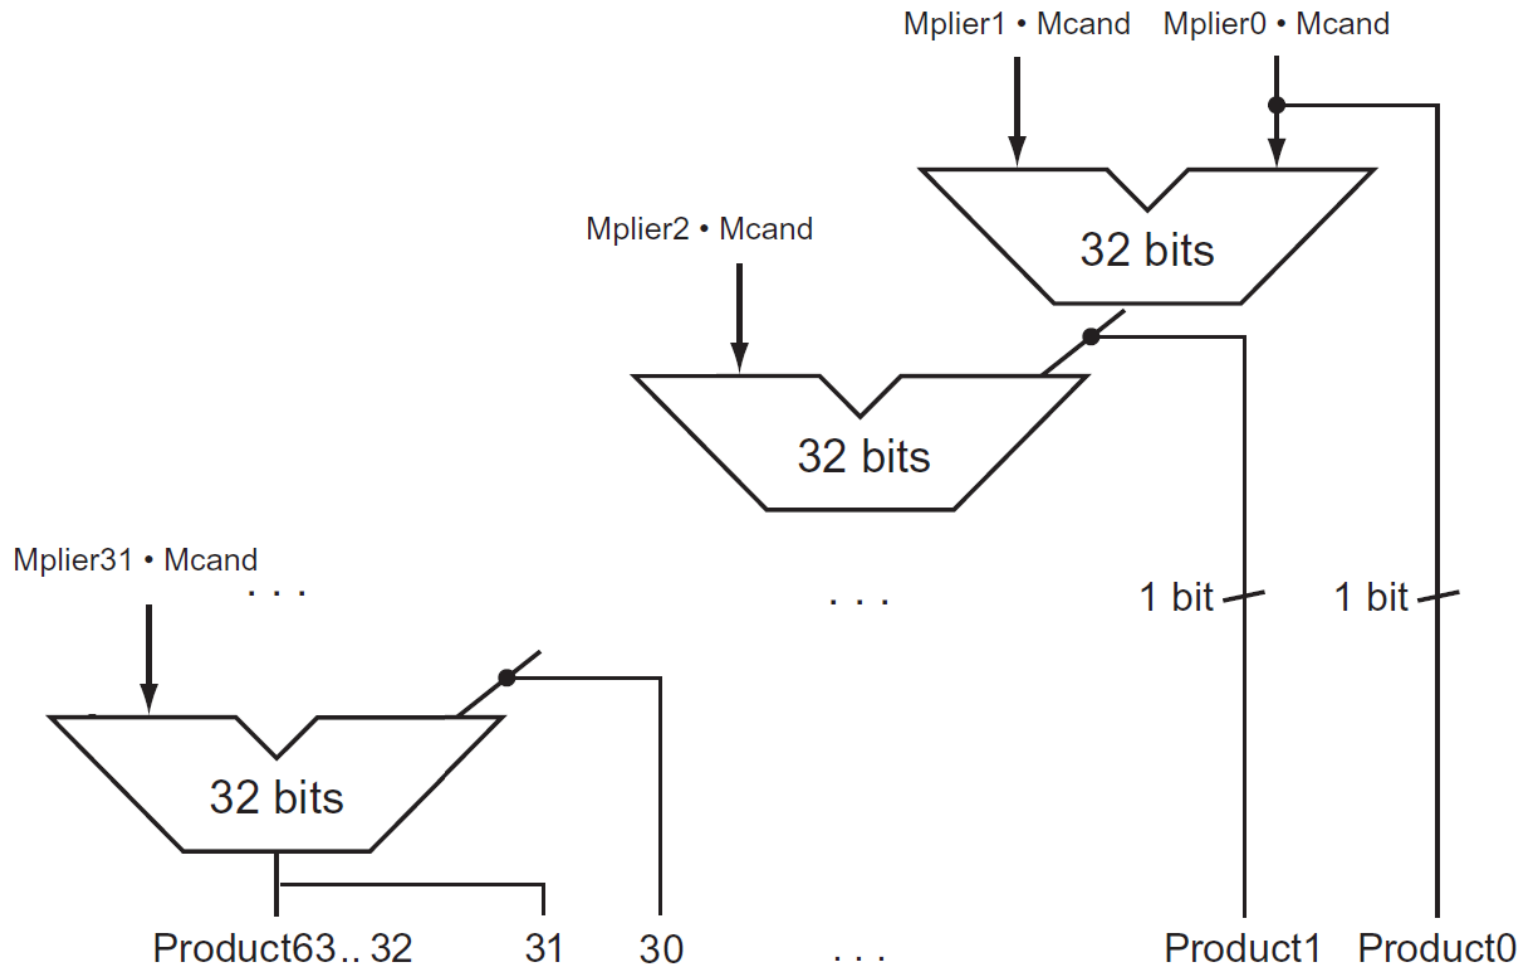
Add and Right Shift Multiplier Hardware



	0 0 0 0	0 1 0 1	= 5
add	0 1 1 0	0 1 0 1	LSB=1 → add multiplicand
	0 0 1 1	→ 0 0 1 0	shift right
add	0 0 1 1	0 0 1 0	LSB=0 → no change
	0 0 0 1	→ 1 0 0 1	shift right
add	0 1 1 1	1 0 0 1	LSB=1 → add multiplicand
	0 0 1 1	→ 1 1 0 0	shift right
add	0 0 1 1	1 1 0 0	LSB=0 → no change
	0 0 0 1	→ 1 1 1 0	shift right
			= 30

Fast multiplier – Design for Moore


❑ Why is this fast?



Fast multiplier – Design for Moore

❑ How fast is this?

❑ Anything wrong?



MIPS Multiply Instruction

- ❑ Multiply (`mult` and `multu`) produces a double precision product (2 x 32 bit)

```
mult    $s0, $s1          # hi||lo = $s0 * $s1
```


0	16	17	0	0	0x18	
---	----	----	---	---	------	--

- ❑ Two additional registers: **hi** and **lo**
- ❑ Low-order word of the product is stored in processor register `lo` and the high-order word is stored in register `hi`
- ❑ Instructions `mfhi rd` and `mflo rd` are provided to move the product to (user accessible) registers in the register file

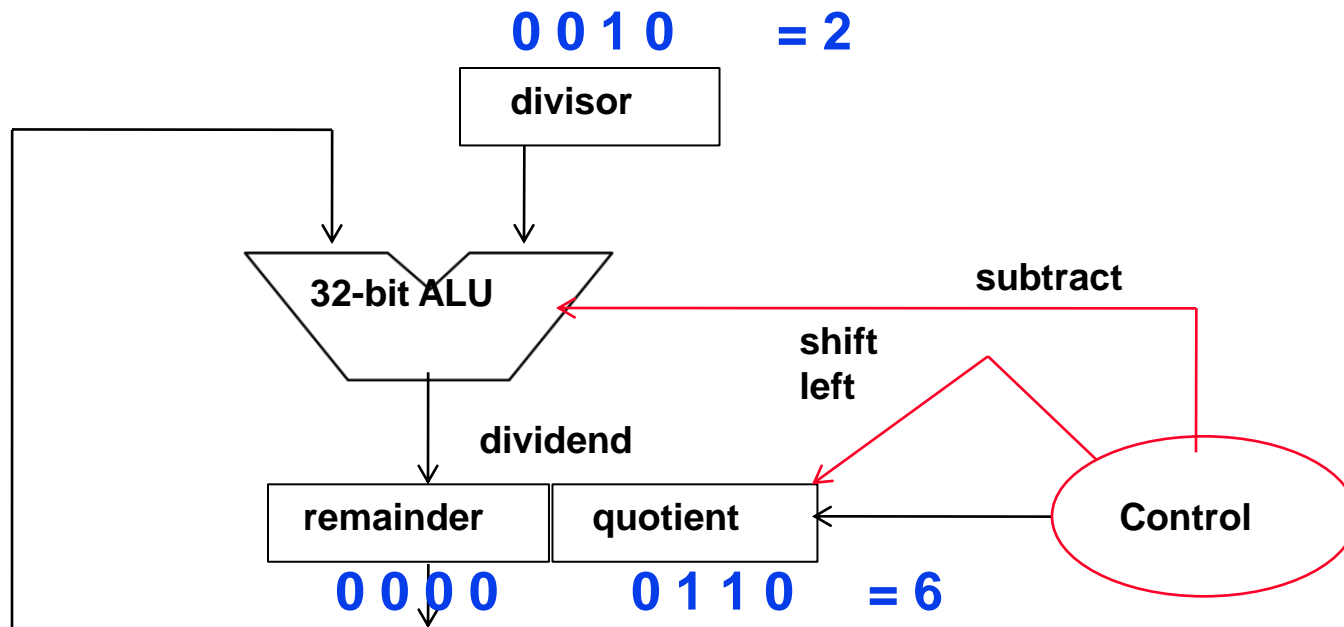
Division

- Division is just a *bunch* of quotient digit guesses and left shifts and subtracts

$$\text{dividend} = \text{quotient} \times \text{divisor} + \text{remainder}$$



Left Shift and Subtract Division Hardware



	0 0 0 0	←	1 1 0 0	
sub	1 1 1 0		1 1 0 0	rem neg, so 'ient bit = 0
	0 0 0 0		1 1 0 0	restore remainder
	0 0 0 1	←	1 0 0 0	
sub	1 1 1 1		1 1 0 0	rem neg, so 'ient bit = 0
	0 0 0 1		1 0 0 0	restore remainder
	0 0 1 1	←	0 0 0 0	
sub	0 0 0 1		0 0 0 1	rem pos, so 'ient bit = 1
	0 0 1 0	←	0 0 1 0	
sub	0 0 0 0		0 0 1 1	rem pos, so 'ient bit = 1
				= 3 with 0 remainder

MIPS Divide Instruction

- Divide (`div` and `divu`) generates the remainder in `hi` and the quotient in `lo`

`div $s0, $s1` # `lo` = `$s0` / `$s1`

 # `hi` = `$s0` mod `$s1`


0	16	17	0	0	0x1A	
---	----	----	---	---	------	--

- Instructions `mfhi rd` and `mflo rd` are provided to move the quotient and remainder to (user accessible) registers in the register file
- As with multiply, divide ignores overflow so software must determine if the quotient is too large. Software must also check the divisor to avoid division by 0.

Signed integer multiplication and division

- ❑ Reuse unsigned multiplication then fix product sign later
- ❑ Multiplication
 - ❑ Multiplicand and multiplier are of the same sign: keep product
 - ❑ Multiplicand and multiplier are of different sign: negate product
- ❑ Division:
 - ❑ Dividend and divisor of the same sign:
 - Keep quotient
 - Keep/negate remainder so it is of the same sign with dividend
 - ❑ Dividend and divisor of different sign:
 - Negate quotient
 - Keep/negate remainder so it is of the same sign with dividend

Signed integer with Booth algorithm



Representing Big (and Small) Numbers

❑ Encoding non-integer value?

- [illegible]

□ PI number

PI = 3.14159....

❑ Problem: how to represent the above numbers?

➔ We need reals or floating-point numbers!

➔ Floating point numbers in decimal:

→ 1000

→ 1×10^3

→ 0.1×10^4

Floating point number

- ❑ In decimal system

$$2013.1228 = 201.31228 * 10$$

$$= 20.131228 * 10^2$$

$$= 2.0131228 * 10^3$$

$$= 20131228 * 10^{-4}$$

- ❑ What is the “standard” form?

$$2.0131228 * 10^3 = \underline{2.0131228} \underline{E+03}$$

mantissa

exponent

- ❑ In binary $X = \pm 1.xxxxx * 2^{yyyy}$


- ❑ ***Sign, mantissa, and exponent need to be represented***

Floating point number

- ❑ Floating point representation in binary

$$(-1)^{\text{sign}} \times 1.F \times 2^{E-\text{bias}}$$

- ❑ Still have to fit everything in 32 bits (single precision)
- ❑ Bias = 127 with single precision floating point number



1 sign bit

8 bits

23 bits

- ❑ Defined by the IEEE 754-1985 standard

- ❑ Single precision: 32 bit

- ❑ Double precision: 64 bit

- ❑ Correspond to float and double in C

Examples

❑ Ex1: convert X into decimal value

$X = 1\textcolor{red}{100}\textcolor{red}{0001}\textcolor{red}{0}101\ 0110\ 0000\ 0000\ 0000\ 0000$

sign = 1 \rightarrow X is negative

E = 1000 0010 = 130

F = 10101100...00

$\rightarrow X = (-1)^1 \times 1.101011000..00 \times 2^{130-127}$

$= -1.101011 \times 2^3 = -1101.011$

$= -13.375$

Example

❑ Ex2: find decimal value of X

X = 0011 1111 1000 0000 0000 0000 0000 0000

sign = 0

e = 0111 1111 = 127

m = 000...0000 (23 bit 0)

$X = (-1)^0 \times 1.00...000 \times 2^{127-127} = 1.0$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

Converting X to plain binary


$$9_{10} = 1001_2$$

$$0.6875 \times 2 = 1.375 \quad \rightarrow \text{get bit } 1$$

$$0.375 \times 2 = 0.75 \quad \rightarrow \text{get bit } 0$$

$$0.75 \times 2 = 1.5 \quad \rightarrow \text{get bit } 1$$

$$0.5 \times 2 = 1.0 \quad \rightarrow \text{get bit } 1$$



$$\rightarrow 9.6875_{10} = 1001.1011_2$$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

$$X = 9.6875_{(10)} = 1001.1011_{(2)} = 1.0011011 \times 2^3$$

Then

$$S = 0$$

$$e = 127 + 3 = 130_{(10)} = 1000\ 0010_{(2)}$$

$$m = 001101100\dots00 \text{ (23 bit)}$$

Finally

$$X = 0100\ 0001\ 0001\ 1011\ 0000\ 0000\ 0000\ 0000$$

Examples

□ $1.0_2 \times 2^{-1} =$


□ $100.75_{10} =$

Some special values


- ❑ Smallest+: 0 00000001 1.00000000000000000000000000000000
= $1 \times 2^{1-127}$
- ❑ Zero: 0 00000000 00000000000000000000000000000000
= true 0
- ❑ Largest+: 0 11111110 1.11111111111111111111111111111111
= $(2-2^{-23}) \times 2^{254-127}$

Too large or too small values

- ❑ **Overflow** (floating point) happens when a positive exponent becomes too large to fit in the exponent field
- ❑ **Underflow** (floating point) happens when a negative exponent becomes too large to fit in the exponent field



- ❑ Reduce the chance of underflow or overflow is to offer another format that has a larger exponent field
- ❑ **Double precision – takes two MIPS words**



Reduce underflow with the same bit length?

- ❑ De-normalized number

IEEE 754 FP Standard Encoding

- ❑ Special encodings are used to represent unusual events
 - ❑ \pm infinity for division by zero
 - ❑ NaN (not a number) for invalid operations such as 0/0
 - ❑ True zero is the bit string all zero

Single Precision		Double Precision		Object Represented
E (8)	F (23)	E (11)	F (52)	
0000 0000	0	0000 ... 0000	0	true zero (0)
0000 0000	nonzero	0000 ... 0000	nonzero	\pm denormalized number
0111 1111 to +127,-126	anything	0111 ...1111 to +1023,-1022	anything	\pm floating point number
1111 1111	+ 0	1111 ... 1111	- 0	\pm infinity
1111 1111	nonzero	1111 ... 1111	nonzero	not a number (NaN)

Floating Point Addition

□ Addition (and subtraction)

$$(\pm F1 \times 2^{E1}) + (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- **Step 0: Restore the hidden bit in F1 and in F2**
- **Step 1: Align fractions by right shifting F2 by $E1 - E2$ positions (assuming $E1 \geq E2$) keeping track of (three of) the bits shifted out in G R and S**
- **Step 2: Add the resulting F2 to F1 to form F3**
- **Step 3: Normalize F3 (so it is in the form 1.XXXXXX ...)**
 - If F1 and F2 have the same sign $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment $E3$ (check for overflow)
 - If F1 and F2 have different signs \rightarrow F3 may require *many* left shifts each time decrementing $E3$ (check for underflow)
- **Step 4: Round F3 and possibly normalize F3 again**
- **Step 5: Rehide the most significant bit of F3 before storing the result**

Floating Point Addition Example

□ Add

$$(0.5 = 1.0000 \times 2^{-1}) + (-0.4375 = -1.1100 \times 2^{-2})$$

□ **Step 0:**

□ **Step 1:**

□ **Step 2:**

□ **Step 3:**

□ **Step 4:**

□ **Step 5:**

Floating Point Addition Example

□ Add: $0.5 + (-0.4375) = ?$

$$(0.5 = 1.0000 \times 2^{-1}) + (-0.4375 = -1.1100 \times 2^{-2})$$

- **Step 0:** Hidden bits restored in the representation above
- **Step 1:** Shift significand with the smaller exponent (1.1100) right until its exponent matches the larger exponent (so once)
- **Step 2:** Add significands
 $1.0000 + (-0.111) = 1.0000 - 0.111 = 0.001$
- **Step 3:** Normalize the sum, checking for exponent over/underflow
 $0.001 \times 2^{-1} = 0.010 \times 2^{-2} = .. = 1.000 \times 2^{-4}$
- **Step 4:** The sum is already rounded, so we're done
- **Step 5:** Re-hide the hidden bit before storing

Floating Point Multiplication

❑ Multiplication

$$(\pm F1 \times 2^{E1}) \times (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- ❑ **Step 0: Restore the hidden bit in F1 and in F2**
- ❑ **Step 1: Add the two (biased) exponents and subtract the bias from the sum, so $E1 + E2 - 127 = E3$**
also determine the sign of the product (which depends on the sign of the operands (most significant bits))
- ❑ **Step 2: Multiply F1 by F2 to form a double precision F3**
- ❑ **Step 3: Normalize F3 (so it is in the form 1.XXXXXX ...)**
 - Since F1 and F2 come in normalized $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment E3
 - Check for overflow/underflow
- ❑ **Step 4: Round F3 and possibly normalize F3 again**
- ❑ **Step 5: Rehide the most significant bit of F3 before storing the result**

Floating Point Multiplication Example

❑ Multiply

$$(0.5 = 1.0000 \times 2^{-1}) \times (-0.4375 = -1.1100 \times 2^{-2})$$

❑ **Step 0:**

❑ **Step 1:**

❑ **Step 2:**

❑ **Step 3:**

❑ **Step 4:**

❑ **Step 5:**

Floating Point Multiplication Example

❑ Multiply

$$(0.5 = 1.0000 \times 2^{-1}) \times (-0.4375 = -1.1100 \times 2^{-2})$$

❑ **Step 0:** Hidden bits restored in the representation above

❑ **Step 1:** Add the exponents (not in bias would be $-1 + (-2) = -3$ and in bias would be $(-1+127) + (-2+127) - 127 = (-1 -2) + (127+127-127) = -3 + 127 = 124$)

❑ **Step 2:** Multiply the significands

$$1.0000 \times 1.110 = 1.110000$$

❑ **Step 3:** Normalized the product, checking for exp over/underflow

$$1.110000 \times 2^{-3} \text{ is already normalized}$$

❑ **Step 4:** The product is already rounded, so we're done

❑ **Step 5:** Rehide the hidden bit before storing

Support for Accurate Arithmetic

- ❑ IEEE 754 FP rounding modes

- ❑ Always round up (toward $+\infty$)
- ❑ Always round down (toward $-\infty$)
- ❑ Truncate
- ❑ **Round to nearest even** (when the Guard || Round || Sticky are 100) – always creates a 0 in the least significant (kept) bit of F

- ❑ Rounding (except for truncation) requires the hardware to include extra F bits during calculations

- ❑ Guard bit – used to provide one F bit when shifting left to normalize a result (e.g., when normalizing F after division or subtraction)
- ❑ Round bit – used to improve rounding accuracy
- ❑ Sticky bit – used to support **Round to nearest even**; is set to a 1 whenever a 1 bit shifts (right) through it (e.g., when aligning F during addition/subtraction)

F = 1 . xxxxxxxxxxxxxxxxxxxxxxxxx **G R S**

Example

❑ Calculate:

$$0.2 \times 5 = ?$$

$$0.333 \times 3 = ?$$

$$(1.0/3) \times 3 = ?$$

MIPS Arithmetic Logic Unit (ALU)

- ❑ Must support the Arithmetic/Logic operations of the ISA

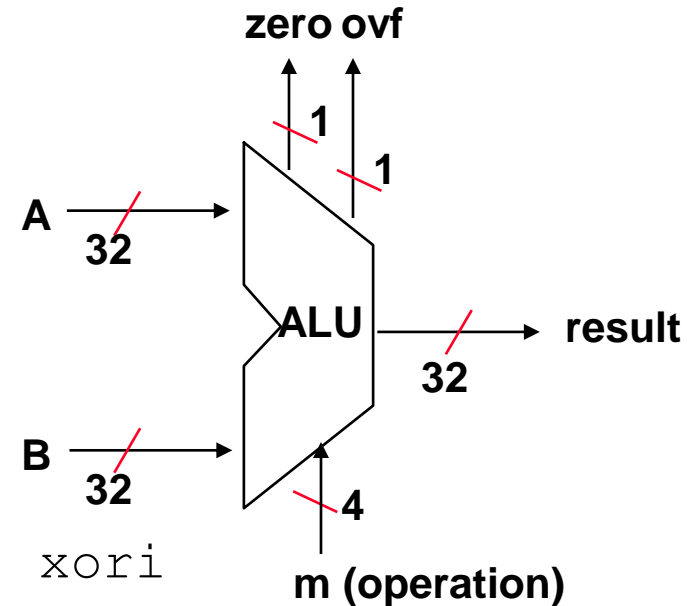
add, addi, addiu, addu

sub, subu

mult, multu, div, divu

and, andi, nor, or, ori, xor, xori

beq, bne, slt, slti, sltiu, sltu



- ❑ With special handling for
 - ❑ **sign extend** – addi, addiu, slti, sltiu
 - ❑ **zero extend** – andi, ori, xori
 - ❑ **overflow detection** – add, addi, sub