
IT4272E-COMPUTER SYSTEMS

Chapter 6: Storage and Other I/O Topics

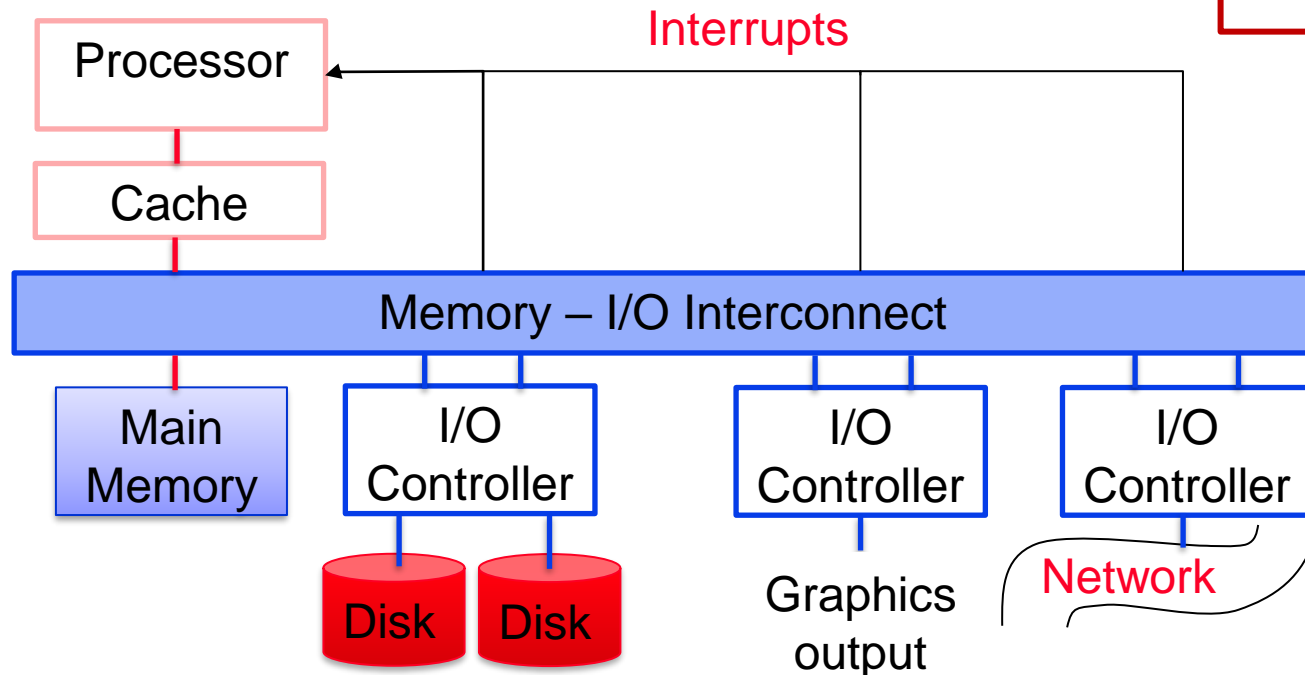


[with materials from *Computer Organization and Design, 4th Edition*,
Patterson & Hennessy, © 2008, MK]

Content

- ❑ 6.1. Introduction
- ❑ 6.2. Dependability, Reliability, and Availability
- ❑ 6.3. Disk Storage
- ❑ 6.4. Flash Storage
- ❑ 6.5. Connecting Processors, Memory, and I/O Devices
- ❑ 6.6. Interfacing I/O Devices
- ❑ 6.7. Measuring I/O Performance
- ❑ 6.8. I/O System Design
- ❑ 6.9 Parallelism and I/O: RAID
- ❑ 6.10. Real Stuff: Server Computer
- ❑ 6.11. Fallacies and Pitfalls
- ❑ 6.12. Concluding Remarks


Introduction

- ❑ I/O devices can be characterized by
 - | Behaviour: input, output, storage
 - | Partner: human or machine
 - | Data rate: bytes/sec, transfers/sec
- ❑ I/O bus connections



I/O System Characteristics

- ❑ Dependability is important
 - | Particularly for storage devices
- ❑ Performance measures
 - | Latency (response time)
 - | Throughput (bandwidth)



Desktops & embedded systems

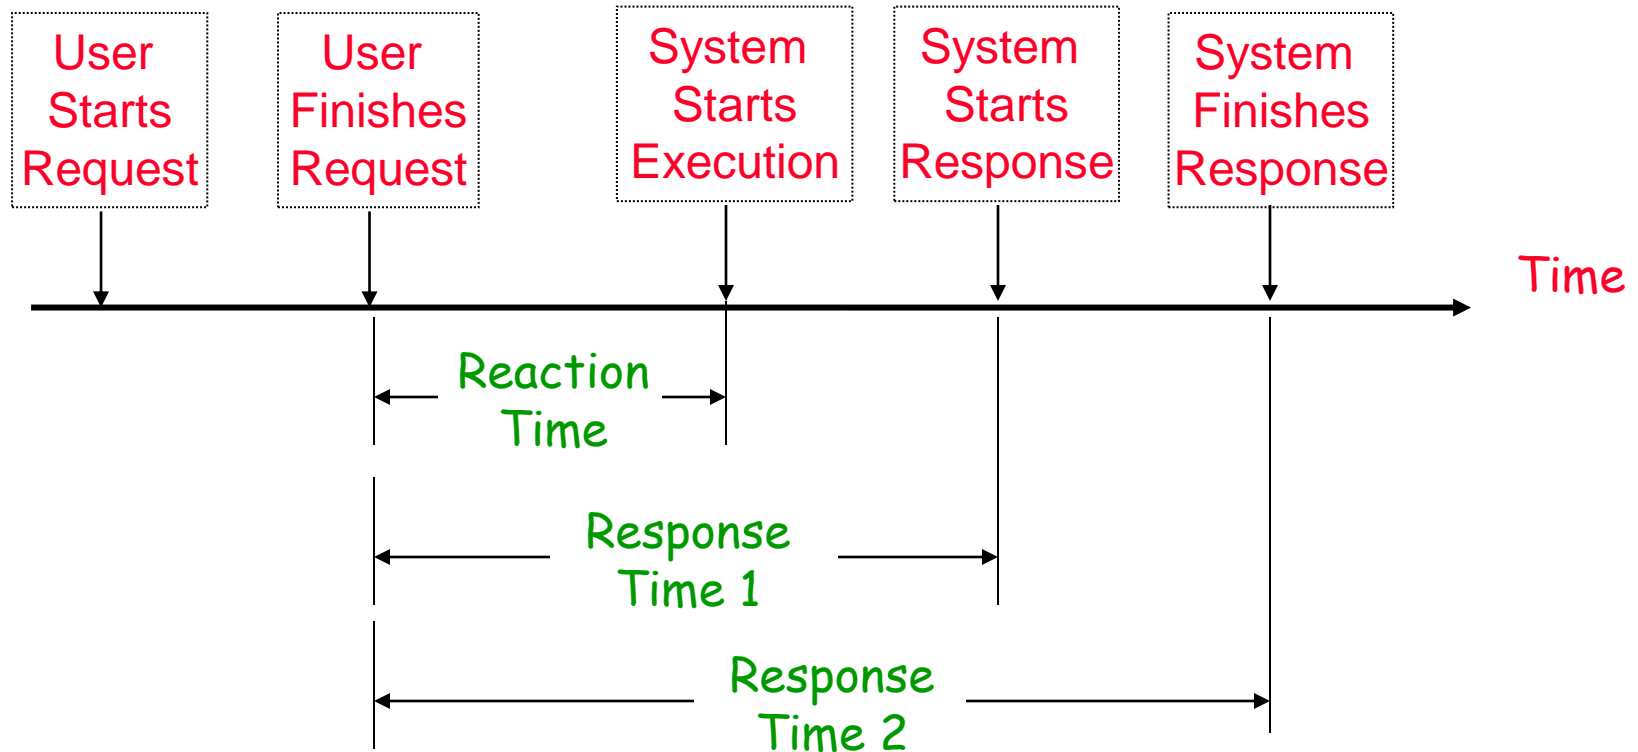
- Mainly interested in: **response time & diversity** of devices



Servers

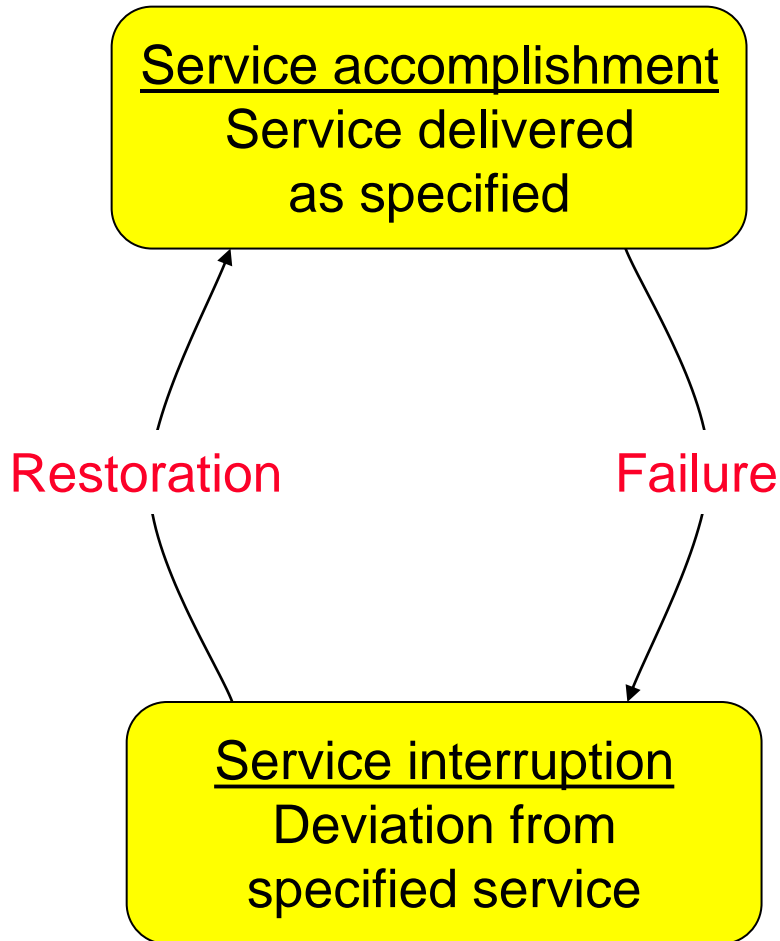
- Mainly interested in: **throughput & expandability** of devices

Respond Time



- Can have two measures of response time
 - Both ok, but 2 preferred if execution long


6.2. Dependability, Reliability, and Availability



- ❑ Fault: failure of a component
 - | May or may not lead to system failure

Dependability Measures

- Reliability: mean time to failure (MTTF)
- Service interruption: mean time to repair (MTTR)
- Mean time between failures
 $MTBF = MTTF + MTTR$
- Availability = $MTTF / (MTTF + MTTR)$



□ Improving Availability




Increase MTTF: fault avoidance, fault tolerance, fault forecasting



Reduce MTTR: improved tools and processes for diagnosis and repair

6.3. Disk Storage

- ❑ Nonvolatile, rotating magnetic storage



Disk Sectors and Access

❑ Each sector records

- | Sector ID
- | Data (512 bytes, 4096 bytes proposed)
- | Error correcting code (ECC)
 - Used to hide defects and recording errors
- | Synchronization fields and gaps

❑ Access to a sector involves

- | Queuing delay if other accesses are pending
- | Seek: move the heads
- | Rotational latency
- | Data transfer
- | Controller **overhead**

Disk Access Example

□ Given

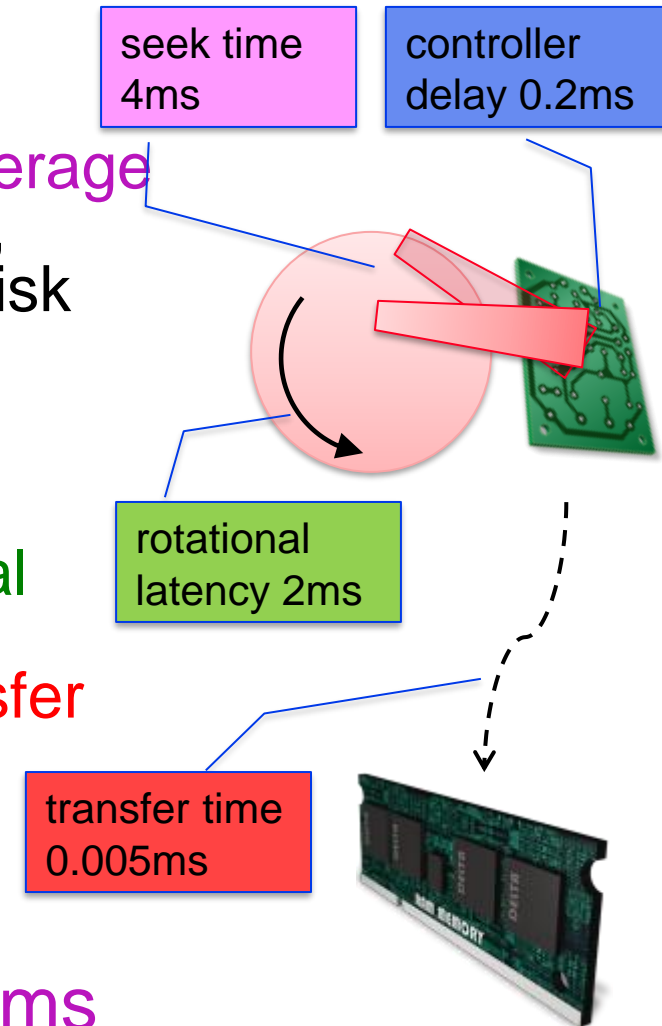
- | 512B sector, 15,000rpm, 4ms average seek time, 100MB/s transfer rate, 0.2ms controller overhead, idle disk

□ Average read time

- | 4ms seek time
+ $\frac{1}{2} / (15,000/60) = 2\text{ms}$ rotational latency
+ $512 / 100\text{MB/s} = 0.005\text{ms}$ transfer time
+ 0.2ms controller delay
= 6.2ms

□ If actual average seek time is 1ms

- | Average read time = 3.2ms




Disk Performance Issues

- ❑ Manufacturers quote average seek time
 - | Based on all possible seeks
 - | Locality and OS scheduling lead to smaller actual average seek times (25%~33%)
- ❑ Smart disk controller allocate physical sectors on disk
 - | Present logical sector interface to host
 - | SCSI, ATA, SATA
- ❑ Disk drives include caches
 - | Prefetch sectors in anticipation of access
 - | Avoid seek and rotational delay

anticipation /æn,tisi'peɪʃn/:
sự đoán trước

6.4. Flash Storage

- ❑ Nonvolatile semiconductor storage
 - | 100× – 1000× faster than disk
 - | Smaller, lower power, more robust
 - | But more \$/GB (between disk and DRAM)



Flash Types

- ❑ NOR flash: bit cell like a NOR gate
 - | Random read/write access
 - | Used for instruction memory in embedded systems
- ❑ NAND flash: bit cell like a NAND gate
 - | Denser (bits/area), but block-at-a-time access
 - | Cheaper per GB
 - | Used for USB keys, media storage, ...
- ❑ Flash bits **wears out** after 1000's of accesses
 - | Not suitable for direct RAM or disk replacement
 - | Wear leveling: remap data to less used blocks

6.5. Connecting Processors, Memory, and I/O Devices

- ❑ Need interconnections between
 - | CPU, memory, I/O controllers
- ❑ Bus: shared communication channel
 - | Parallel set of wires for data and synchronization of data transfer
 - | Can become a bottleneck
- ❑ Performance limited by physical factors
 - | Wire length, number of connections
- ❑ More recent alternative: high-speed serial connections with switches
 - | Like networks

Bus Types

❑ Processor-Memory buses

- | Short, high speed
- | Design is matched to memory organization

❑ I/O buses

- | Longer, allowing multiple connections
- | Specified by standards for interoperability
- | Connect to processor-memory bus through a bridge

Bus Signals and Synchronization

❑ Data lines

- | Carry **address** and data
- | Multiplexed or separate

❑ Control lines

- | Indicate data type, synchronize transactions

❑ Synchronous

- | Uses a bus clock

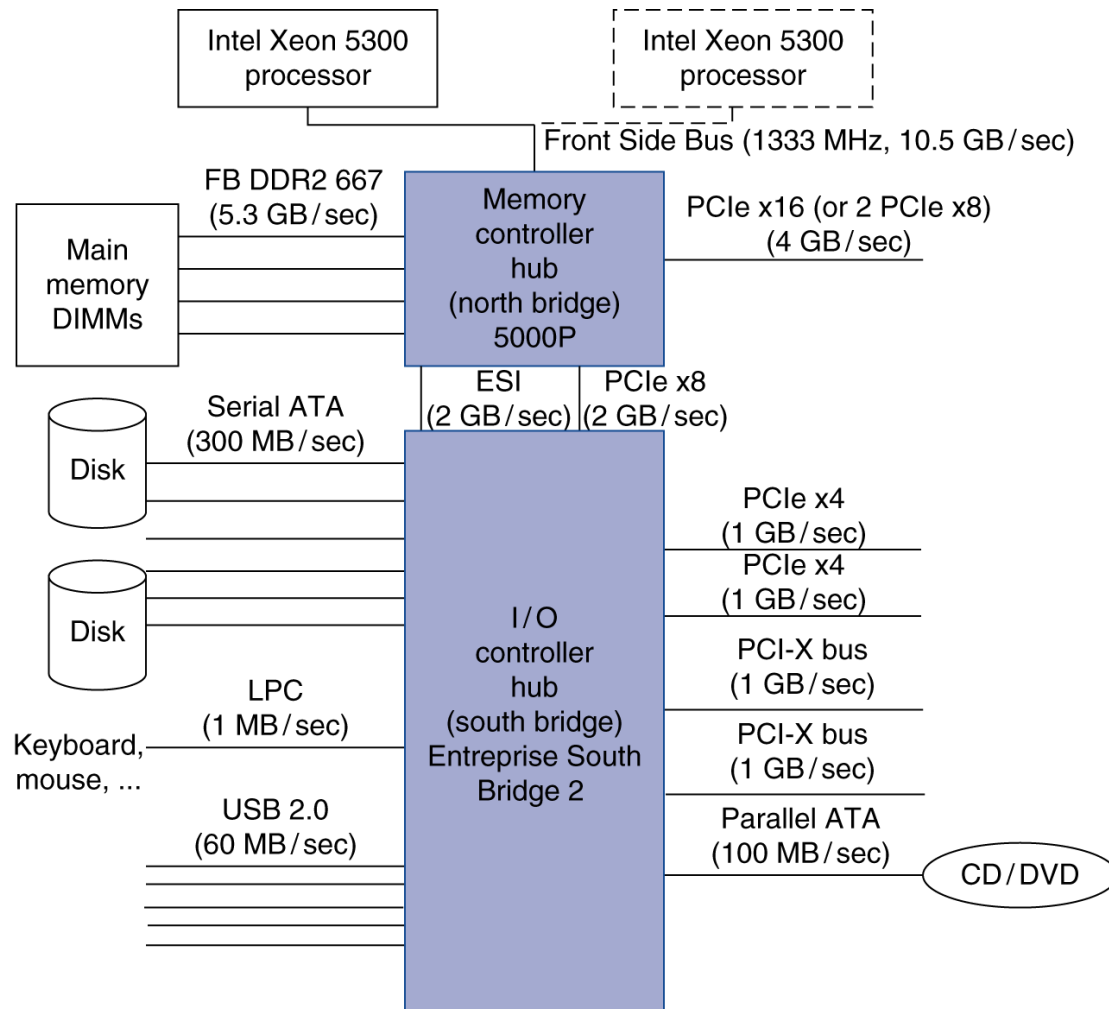
❑ Asynchronous

- | Uses request/acknowledge control lines for handshaking

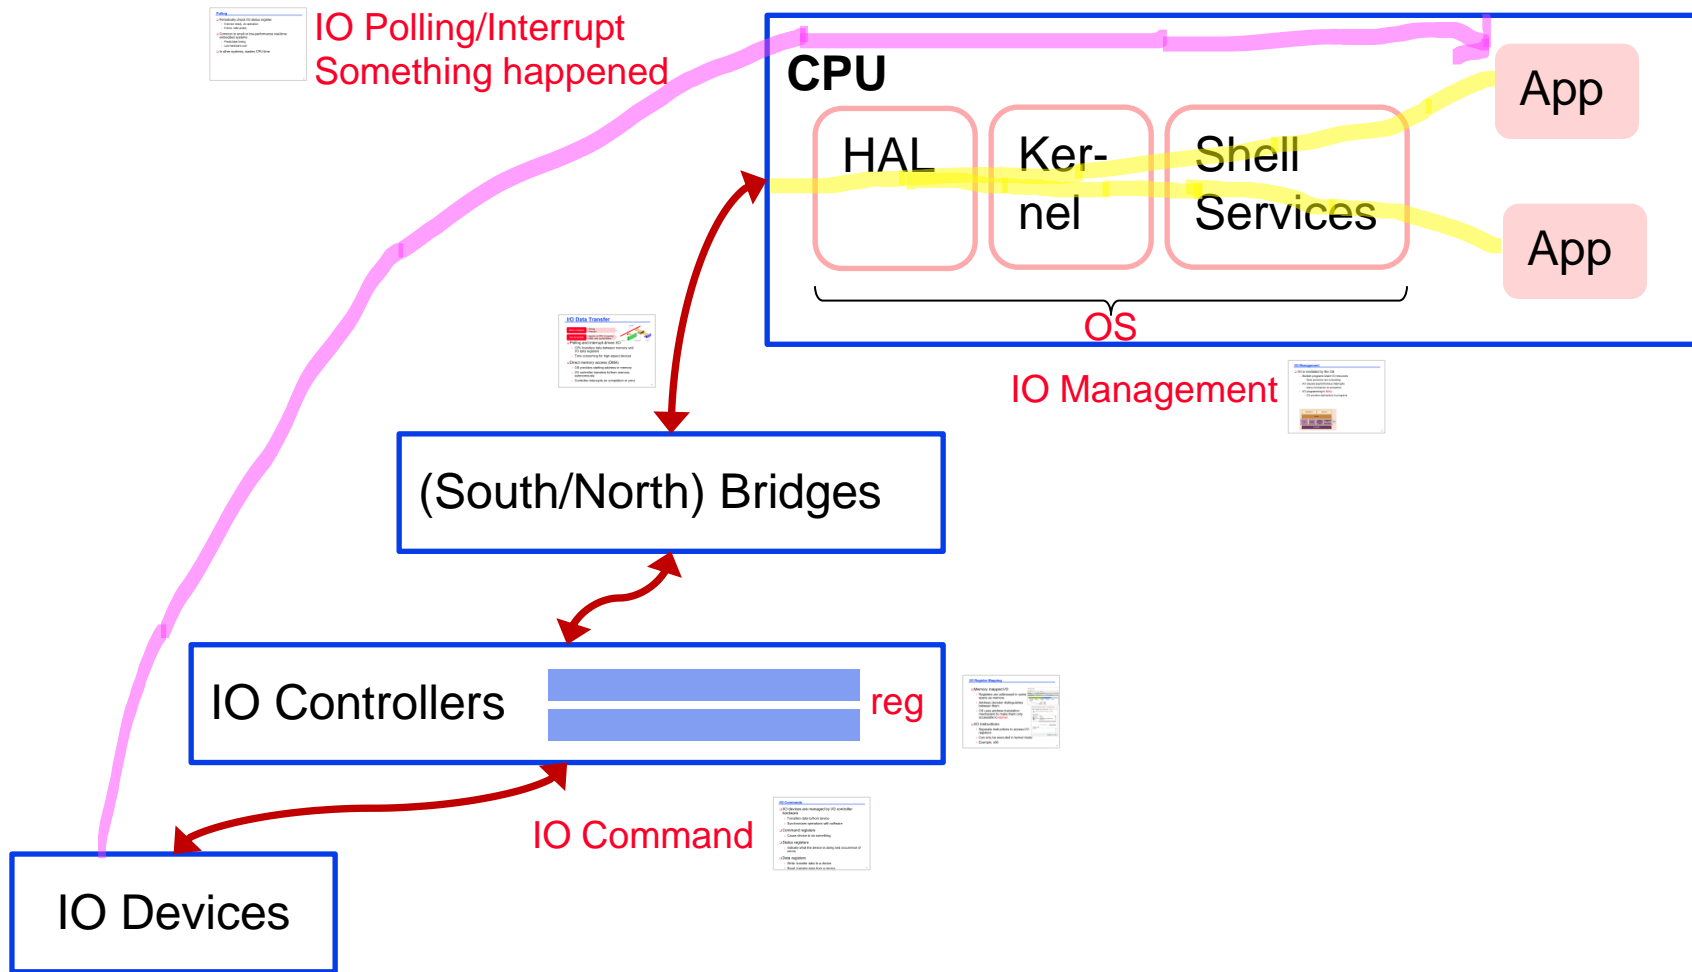
I/O Bus Examples

| | Firewire | USB 2.0 | PCI Express | Serial ATA | Serial Attached SCSI |
|---------------------|-------------------|-----------------------------|--|------------|----------------------|
| Intended use | External | External | Internal | Internal | External |
| Devices per channel | 63 | 127 | 1 | 1 | 4 |
| Data width | 4 | 2 | 2/lane | 4 | 4 |
| Peak bandwidth | 50MB/s or 100MB/s | 0.2MB/s, 1.5MB/s, or 60MB/s | 250MB/s/lane 1×, 2×, 4×, 8×, 16×, 32× | 300MB/s | 300MB/s |
| Hot pluggable | Yes | Yes | Depends | Yes | Yes |
| Max length | 4.5m | 5m | 0.5m | 1m | 8m |
| Standard | IEEE 1394 | USB Implementers Forum | PCI-SIG | SATA-IO | INCITS TC T10 |

Typical x86 PC I/O System

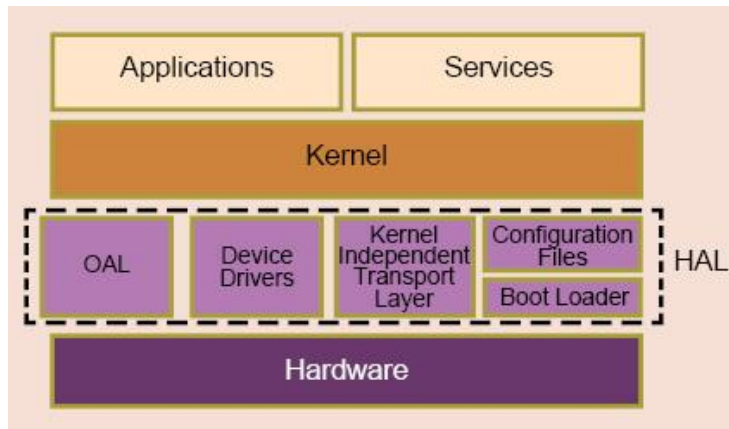


6.6. Interfacing I/O Devices



I/O Management

- ❑ I/O is mediated by the OS
 - | Multiple programs share I/O resources
 - Need protection and scheduling
 - | I/O causes asynchronous interrupts
 - Same mechanism as exceptions
 - | I/O programming is **fiddly**
 - OS provides abstractions to programs




I/O Commands

- ❑ I/O devices are managed by I/O controller hardware
 - | Transfers data to/from device
 - | Synchronizes operations with software
- ❑ Command registers
 - | Cause device to do something
- ❑ Status registers
 - | Indicate what the device is doing and occurrence of errors
- ❑ Data registers
 - | Write: transfer data to a device
 - | Read: transfer data from a device

I/O Register Mapping

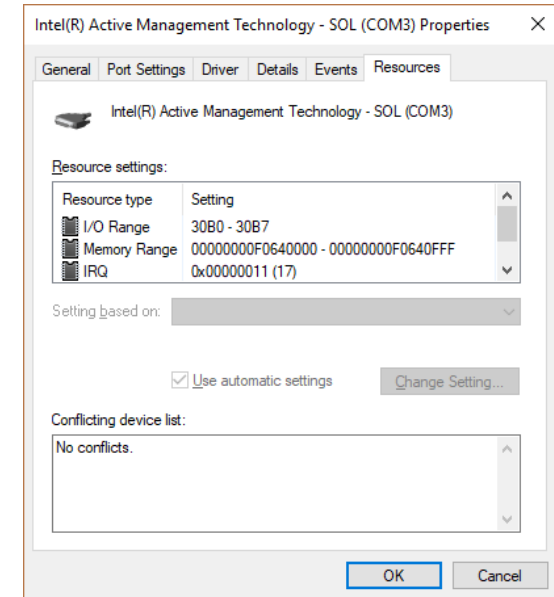
❑ Memory mapped I/O

- | Registers are addressed in same space as memory
- | Address decoder distinguishes between them
- | OS uses address translation mechanism to make them only accessible to **kernel**



❑ I/O instructions

- | Separate instructions to access I/O registers
- | Can only be executed in kernel mode
- | Example: x86



Polling

- ❑ Periodically check I/O status register
 - | If device ready, do operation
 - | If error, take action
- ❑ Common in small or low-performance real-time embedded systems
 - | Predictable timing
 - | Low hardware cost
- ❑ In other systems, wastes CPU time

Interrupts

- ❑ When a device is ready or error occurs
 - | Controller interrupts CPU
- ❑ Interrupt is like an exception
 - | But not synchronized to instruction execution
 - | Can invoke handler between instructions
 - | Cause information often identifies the interrupting device
- ❑ Priority interrupts
 - | Devices needing more urgent attention get higher priority
 - | Can interrupt handler for a lower priority interrupt

Interrupts: Examples

```
tiennnd@tiennnd:~$ cat /proc/interrupts
```

| | CPU0 | CPU1 | CPU2 | CPU3 | | |
|------|---------|---------|--------|--------|-------------------------|-----------------------|
| 0: | 76 | 0 | 0 | 0 | IO-APIC-edge | timer |
| 1: | 13090 | 3 | 0 | 0 | IO-APIC-edge | i8042 |
| 8: | 1 | 0 | 0 | 0 | IO-APIC-edge | rtc0 |
| 9: | 1064 | 0 | 0 | 0 | IO-APIC-fastestoi | acpi |
| 12: | 213112 | 2 | 0 | 0 | IO-APIC-edge | i8042 |
| 16: | 915415 | 0 | 0 | 0 | IO-APIC-fastestoi | ehci_hcd:usb1, nvidia |
| 17: | 282256 | 347 | 0 | 22193 | IO-APIC-fastestoi | ath9k, snd_hda_intel |
| 23: | 74438 | 0 | 0 | 0 | IO-APIC-fastestoi | ehci_hcd:usb2 |
| 41: | 81104 | 0 | 0 | 0 | PCI-MSI-edge | ahci |
| 42: | 0 | 0 | 0 | 0 | PCI-MSI-edge | eth0 |
| 43: | 10 | 0 | 0 | 0 | PCI-MSI-edge | mei |
| 44: | 286 | 0 | 0 | 0 | PCI-MSI-edge | snd_hda_intel |
| NMI: | 39 | 609 | 372 | 359 | Non-maskable interrupts | |
| LOC: | 1377006 | 1000228 | 742879 | 758045 | Local timer interrupts | |
| SPU: | 0 | 0 | 0 | 0 | Spurious interrupts | |

IRQ Number

the number of interrupt handled by CPU Core

Interrupt Type

Device Name

❑ Example with Asus K43SJ

❑ Each CPU in the system has its own column and its own number of interrupts per IRQ.

❑ IRQ0: system timer; IRQ1&12: keyboard&mouse.


I/O Data Transfer

When it happen

- Polling
- Interrupt

How to transfer

- mem/io → CPU → mem/io
- DMA, with cache/VMem



❑ Polling and interrupt-driven I/O

- | CPU transfers data between memory and I/O data registers
- | Time consuming for high-speed devices

❑ Direct memory access (DMA)

- | OS provides starting address in memory
- | I/O controller transfers to/from memory autonomously
- | Controller interrupts on completion or error

DMA/Cache Interaction


- ❑ If DMA writes to a memory block that is cached
 - | Cached copy becomes stale
- ❑ If write-back cache has dirty block, and DMA reads memory block
 - | Reads stale data
- ❑ Need to ensure cache coherence
 - | Flush blocks from cache if they will be used for DMA
 - | Or use non-cacheable memory locations for I/O

stale /steil/ (adj): cũ rích, mất hiệu lực
coherence /kou'hiərəns/: tính nhất quán

DMA/VM Interaction

❑ OS uses virtual addresses for memory

- | DMA blocks may not be contiguous in physical memory



❑ Should DMA use virtual addresses?

- | Would require controller to do translation

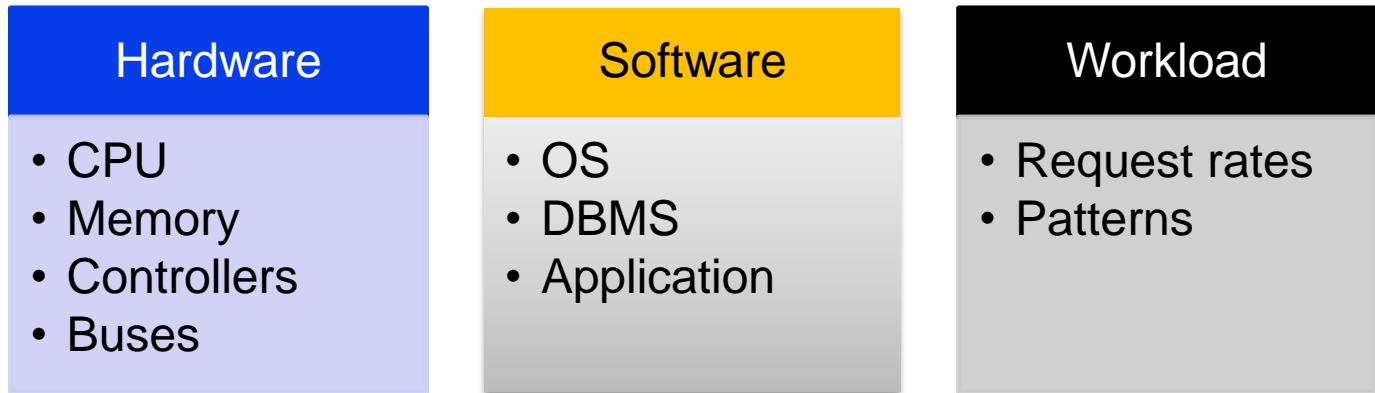
❑ If DMA uses physical addresses

- | May need to break transfers into page-sized chunks
- | Or chain multiple transfers
- | Or allocate contiguous physical pages for DMA

contiguous /kən'tɪɡjuəs/: liền kề, bên cạnh

6.7. Measuring I/O Performance

- ❑ I/O performance depends on




- ❑ I/O system design can trade-off between **response time** and **throughput**
 - ▮ Measurements of **throughput** often done with constrained **response-time**




Throughput vs Response Time

- *Throughput* increases as load increases, to a point



- **Nominal capacity** is ideal (ex: 10 Mbps)
- **Usable capacity** is achievable (ex: 9.8 Mbps)
- **Knee** is where response time goes up rapidly for small increase in throughput



Transaction Processing Benchmarks

❑ Transactions

- | Small data accesses to a DBMS
- | **Interested in I/O rate**, not data rate

❑ Measure throughput

- | Subject to response time limits and failure handling
- | ACID (Atomicity, Consistency, Isolation, Durability)
- | Overall cost per transaction

❑ Transaction Processing Council (TPC) benchmarks (www.tpc.org)

- | TPC-APP: B2B application server and web services
- | TCP-C: on-line order entry environment
- | TCP-E: on-line transaction processing for brokerage firm
- | TPC-H: decision support — business oriented ad-hoc queries

File System & Web Benchmarks

❑ SPEC System File System (SFS)

- | Synthetic workload for NFS server, based on monitoring real systems
- | Results
 - Throughput (operations/sec)
 - Response time (average ms/operation)

❑ SPEC Web Server benchmark

- | Measures simultaneous user sessions, subject to required throughput/session
- | Three workloads: Banking, Ecommerce, and Support

I/O vs. CPU Performance

❑ Amdahl's Law

- | Don't neglect I/O performance as parallelism increases compute performance

❑ Example

- | Benchmark takes 90s CPU time, 10s I/O time
- | Double the number of CPUs/2 years
 - I/O unchanged

| Year | CPU time | I/O time | Elapsed time | % I/O time |
|------|----------|----------|--------------|------------|
| now | 90s | 10s | 100s | 10% |
| +2 | 45s | 10s | 55s | 18% |
| +4 | 23s | 10s | 33s | 31% |
| +6 | 11s | 10s | 21s | 47% |

Amdahl and Gustafson's Laws

- ❑ **Amdahl's Law:** The speed up achieved through parallelization of **a program** is **limited by the percentage of its workload** that is inherently **serial**.

$$\text{Speedup}(N) = 1 / (S + (1-S)/N) < 1/S$$

N: processors, S: proportion of none-parallelization

- ❑ **Gustafson's Law:** With **increasing data size**, the **speedup** obtained through parallelization **increases**, because the parallel work increases with data size.

$$\text{Speedup}(N) = N - S(N-1) = N(1-S) + S$$

denominator /di'nɒmineitə/ mẫu số; mẫu thức

numerator /'nju:məreitə/: tử số

fraction /'frækʃn/: phân số

Amdahl and Gustafson's Laws

❑ ĐỊNH LUẬT AMDAHL và GUSTAFSON

- Amdahl's Law: Sự tăng tốc của một chương trình sử dụng nhiều bộ xử lý trong tính toán song song bị hạn chế bởi một phần tính toán tuần tự của chương trình (dựa trên giả thiết rằng kích thước của bài toán không đổi khi tính toán song song)
- Khi tăng khối lượng dữ liệu cần xử lý, hệ số tăng tốc sẽ tăng lên khi tiến hành song song hóa, bởi vì khối lượng công việc có thể tính toán song song cũng tăng lên theo dữ liệu

Exercise

- ❑ Assume accelerating a machine by adding a vector mode to it. When a computation is run in vector mode, it is 20 times faster than the normal mode of execution.

However, the software program cannot be parallized absolutely and CPU's speedup of this program is only 2. So how many percent the software cannot run in vector mode?


Answer:

6.8. I/O System Design

❑ Satisfying latency requirements

- | For **time-critical** operations
- | If system is unloaded
 - Add up latency of components


❑ Maximizing throughput



- | Find “weakest link” (lowest-bandwidth component)
 - | Configure to operate at its maximum bandwidth
 - | Balance remaining components in the system
- ### ❑ If system is loaded, simple analysis is insufficient
- | Need to use queuing models or simulation

6.9 Parallelism and I/O: RAID

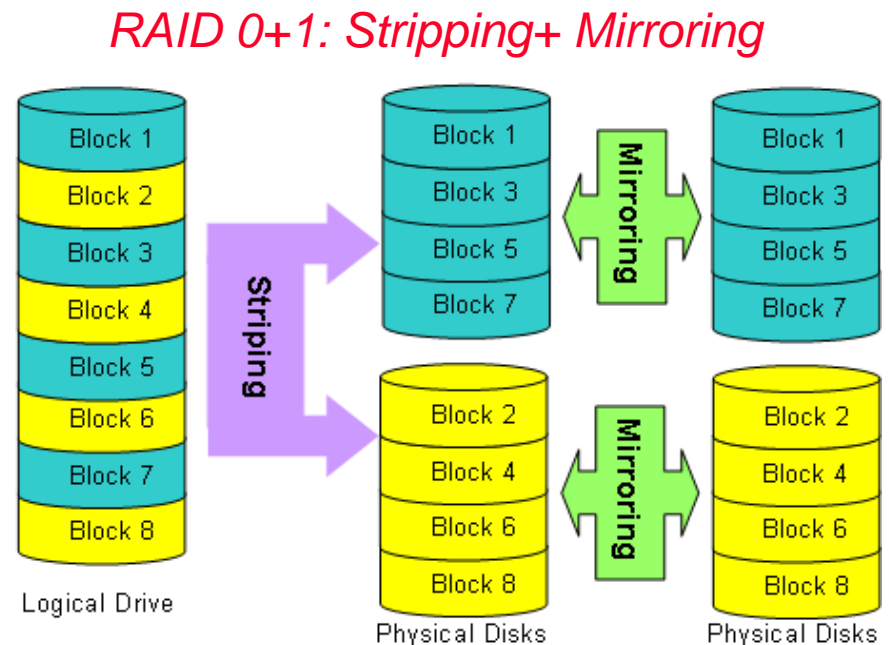

- ❑ Redundant Array of Inexpensive (Independent) Disks
 - | Use multiple smaller disks (c.f. one large disk)
 - | Parallelism improves performance
 - | Plus extra disk(s) for redundant data storage
- ❑ Provides fault tolerant storage system
 - | Especially if failed disks can be “hot swapped”
- ❑ RAID 0, stripping
 - | No redundancy (“AID”?)
 - Just stripe data over multiple disks
 - | But it does improve performance



RAID 1 & 0+1

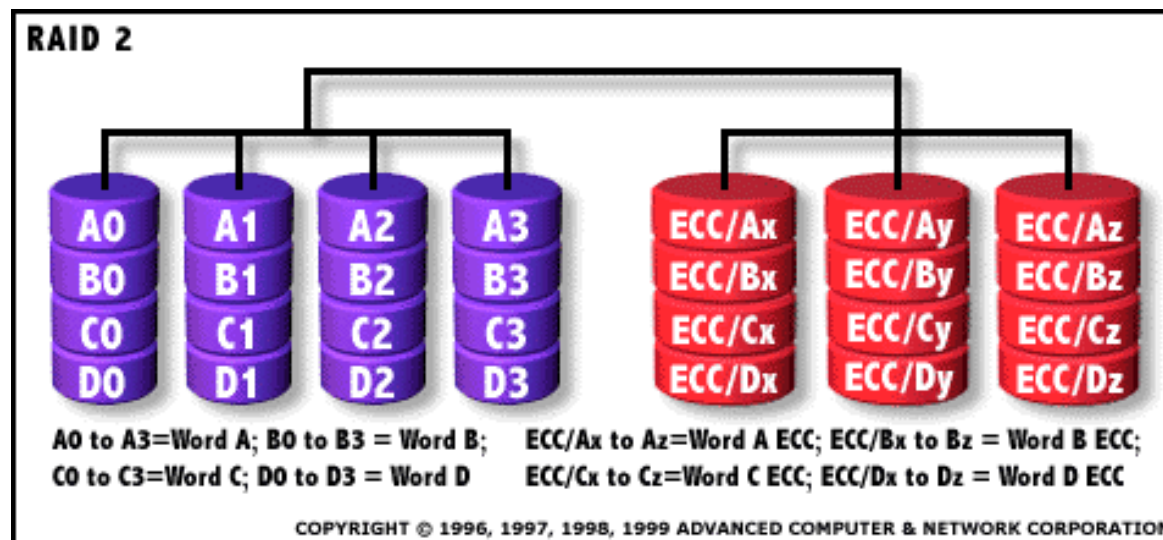
❑ RAID 1: Mirroring

- | N + N disks, replicate data
 - Write data to both data disk and mirror disk
 - On disk failure, read from mirror



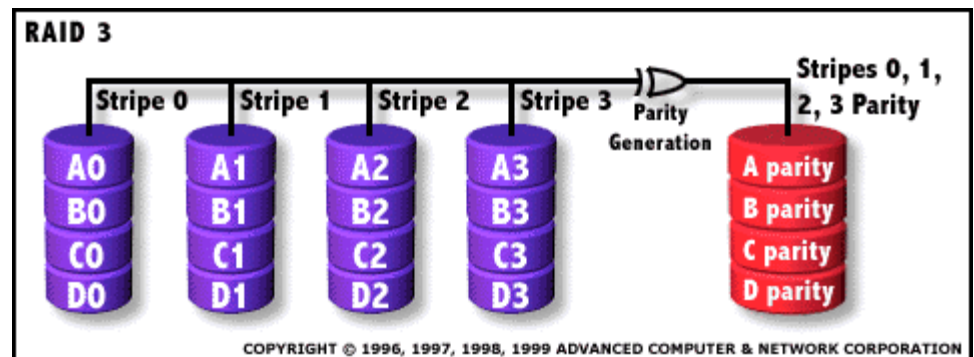
RAID 2, bit striped

- ❑ RAID 2: Error correcting code (ECC)
 - | N + E disks (e.g., 10 + 4)
 - | Split data at **bit level** across N disks
 - | Generate E-bit ECC
 - | Too complex, not used in practice



RAID 3: Bit-Interleaved Parity

- ❑ $N + 1$ disks
 - | Data striped across N disks at **byte level**
 - | Redundant disk stores parity (dedicated parity disk)
 - | **Read access:** Read all disks
 - | **Write access:** Generate new parity and update all disks
 - | **On failure:** Use parity to reconstruct missing data
- ❑ Not widely used

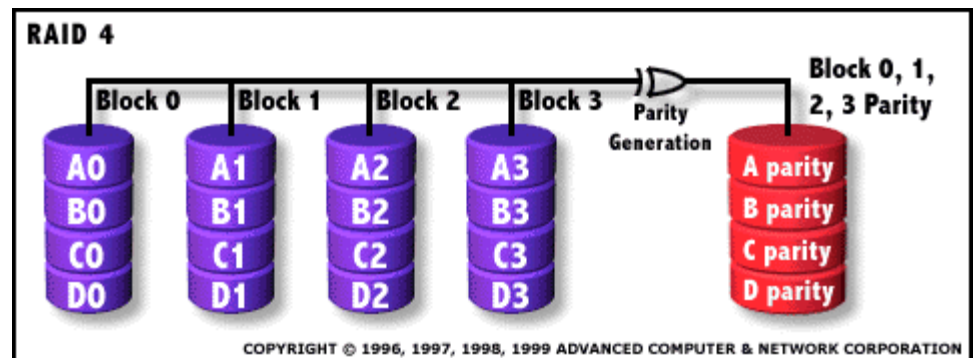


RAID 4: Block-Interleaved Parity

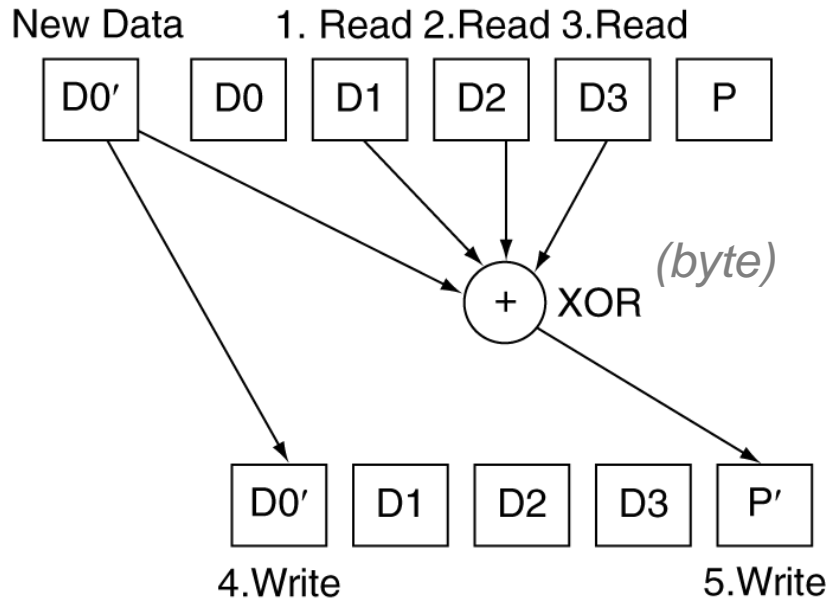
❑ N + 1 disks

- | Data striped across N disks at **block level** (16, 32, 64, 128 kB)
- | Redundant disk stores parity for a group of blocks
- | Read access
 - Read only the disk holding the required block
- | Write access
 - Just read disk containing modified block, and parity disk
 - Calculate new parity, update data disk and parity disk
- | On failure
 - Use parity to reconstruct missing data

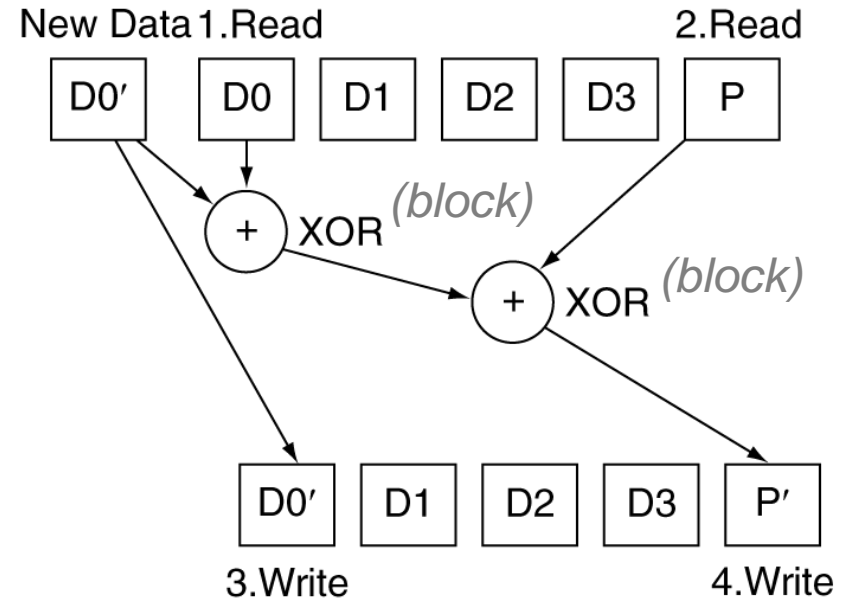
❑ Not widely used



RAID 3 vs RAID 4



*Read 3 disks to get 3 bytes,
and then create parity byte*




*Read 2 disks to get 2 blocks
(include parity block), and then
create new parity block*

RAID 5: Distributed Parity

- ❑ $N + 1$ disks

- | Like RAID 4, but parity blocks distributed across disks
 - Avoids parity disk being a bottleneck

- ❑ Widely used




RAID 6: P + Q Redundancy

❑ N + 2 disks

- | Like RAID 5, but two lots of parity
- | Greater fault tolerance through more redundancy

❑ Multiple RAID

- | More advanced systems give similar fault tolerance with better performance



RAID Summary

- ❑ RAID can improve **performance** and **availability**
 - | High **availability** requires hot swapping
- ❑ Assumes independent disk failures
 - | Too bad if the building burns down!
- ❑ See “Hard Disk Performance, Quality and Reliability”
 - | <http://www.pcguide.com/ref/hdd/perf/index.htm>



6.10. Real Stuff: Server Computer

- ❑ Applications are increasingly run on servers
 - | Web search, office apps, virtual worlds, cloud...
- ❑ Requires large data centre servers
 - | Multiple processors, networks connections, massive storage
 - | Space and power constraints
- ❑ Server equipment built for 19" racks
 - | Multiples of 1.75" (1U) high

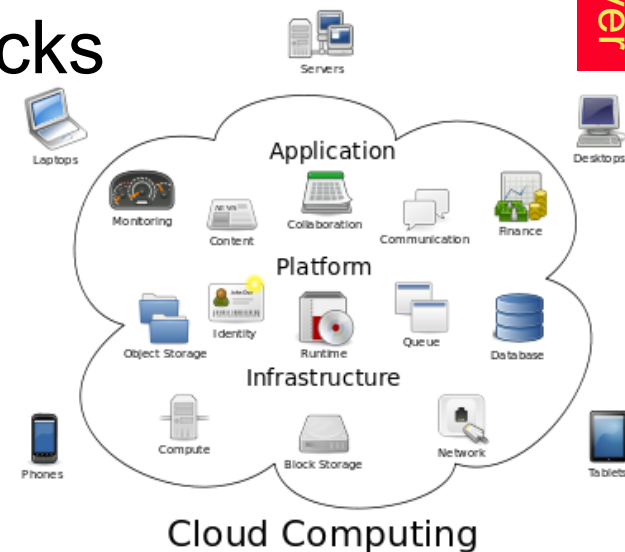


1U = 1.75"



19"


2U



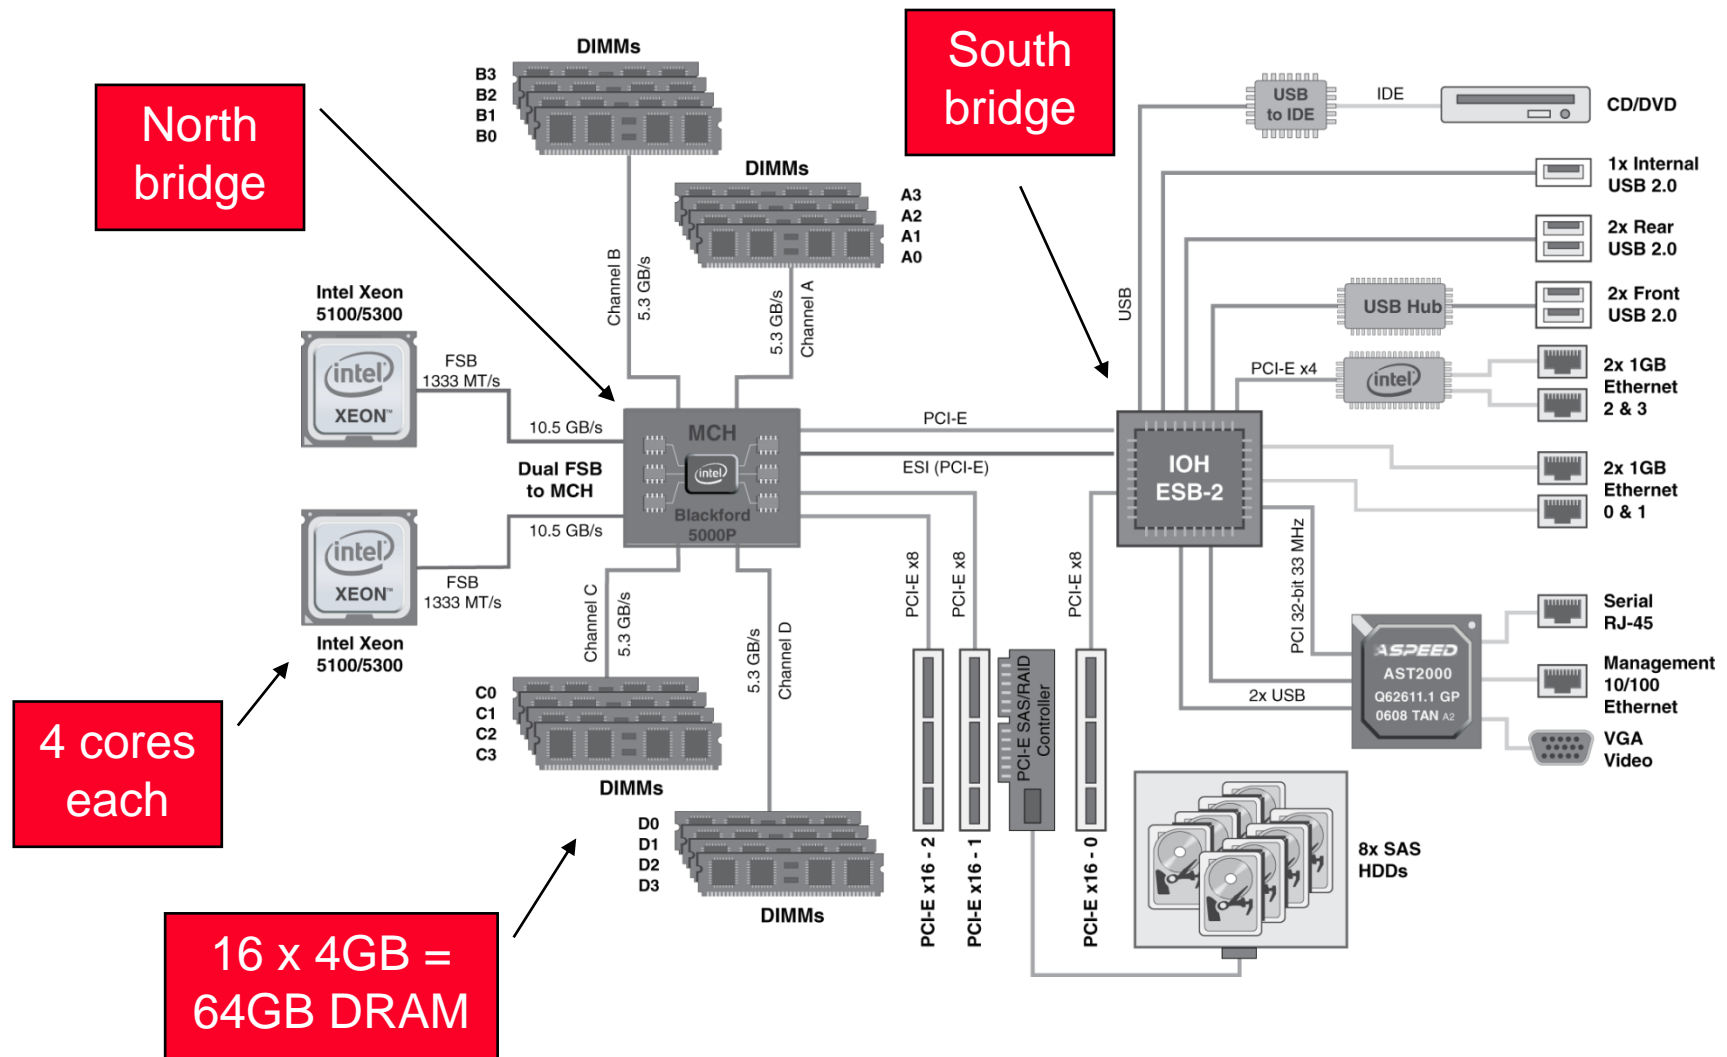
Rack-Mounted Servers



Sun Fire x4150 1U server



Sun Fire x4150 1U server



I/O System Design Example

- ❑ Given a Sun Fire x4150 system with
 - | Workload: 64KB disk reads
 - Each I/O op requires 200,000 user-code instructions and 100,000 OS instructions
 - | Each CPU: 10^9 instructions/sec
 - | FSB: 10.6 GB/sec peak
 - | DRAM DDR2 667MHz: 5.336 GB/sec
 - | PCI-E 8x bus: $8 \times 250\text{MB/sec} = 2\text{GB/sec}$
 - | Disks: 15,000 rpm, 2.9ms avg. seek time, 112MB/sec transfer rate
- ❑ What I/O rate can be sustained?
 - | For random reads, and for sequential reads

Design Example (cont)

❑ I/O rate for CPUs

- Per core: $10^9 / (100,000 + 200,000) = 3,333$ IOs/sec
- 8 cores: $3,333 \times 8 = 26,667$ IOs/sec

- ❑ Random reads, I/O rate for disks

- Assume actual seek time is average/4
- Time/op = seek + latency + transfer (~~+ control time = 0~~)

$$= 2.9\text{ms}/4 + 4\text{ms}/2 + 64\text{KB}/(112\text{MB/s}) = 3.3\text{ms}(\text{per IOs})$$
- $1000 * 1/3.3 = 303$ IOs/sec per disk, 2424 IOs/sec for 8 disks



❑ Sequential reads

- | 112MB/s / 64KB = 1750 IOs/sec per disk
- | 14,000 ops/sec for 8 disks

Design Example (cont)

❑ PCI-E I/O rate

- | $2\text{GB/sec} / 64\text{KB} = 31,250 \text{ IOs/sec}$

❑ DRAM I/O rate

- | $5.336 \text{ GB/sec} / 64\text{KB} = 83,375 \text{ IOs/sec}$

❑ FSB I/O rate

- | Assume we can sustain half the peak rate

- | $10.6 \text{ GB/sec} / 2 / 64\text{KB} = 81,540 \text{ IOs/sec per FSB}$

- | 163,080 IOs/sec for 2 FSBs (2 Intel Xeon)

❑ Weakest link: disks

- | 2424 IOs/sec random, 14,000 IOs/sec sequential

- | Other components have ample headroom to accommodate these rates

*ample /'æmpl/ (adj): nhiều, phong phú
headroom: không gian trống*

6.11. Fallacy: Disk Dependability

- ❑ If a disk manufacturer quotes MTTF as 1,200,000hr (140yr)
 - | A disk will work that long
- ❑ Wrong: this is the mean time to failure
 - | What is the distribution of failures?
 - | What if you have 1000 disks
 - How many will fail per year?



$$\text{Failed Disks} = \frac{1000 \text{ disks} \times (24 * 365) \text{ hrs/disk}}{1200000 \text{ hrs/failure}} = 7.3$$

$$\text{Annual Failure Rate (AFR)} = \frac{7.3}{1000} = 0.73\%$$

fallacy /'fæləsi/ ảo tưởng; ý kiến sai lầm

Fallacies



Prof. Bianca Schroeder

❑ Disk failure rates are as specified

- | Studies of failure rates in the field
 - Schroeder and Gibson: 2% to 4% vs. 0.6% to 0.8%
 - Pinheiro, *et al.*: 1.7% (first year) to 8.6% (third year) vs. 1.5%
- | Why?

❑ A 1GB/s interconnect transfers 1GB in one sec

- | But what's a GB?
- | For bandwidth, use $1\text{GB} = 10^9 \text{ B}$
- | For storage, use $1\text{GB} = 2^{30} \text{ B} = 1.075 \times 10^9 \text{ B}$
- | So 1GB/sec is 0.93GB in one second
 - About 7% error



Pitfall: Offloading to I/O Processors

- ❑ Overhead of managing I/O processor request may dominate
 - | Quicker to do small operation on the CPU
 - | But I/O architecture may prevent that
- ❑ I/O processor may be slower
 - | Since it's supposed to be simpler
- ❑ Making it faster makes it into a major system component
 - | Might need its own coprocessors!



pitfall /'pitfɔ:l/ cạm bẫy
Offload: đẩy dữ liệu ra ngoài vi

Pitfall: Backing Up to Tape

- ❑ Magnetic tape used to have advantages
 - | Removable, high capacity
- ❑ Advantages eroded by disk technology developments
- ❑ Makes better sense to replicate data
 - | E.g, RAID, remote mirroring



Tape

IBM System
Storage TS1130
Tape Drive



erode /i'roud/: xói mòn, suy giảm

Fallacy: Disk Scheduling

- ❑ Best to let the OS schedule disk accesses
 - | But modern drives deal with **L**ogical **B**lock **A**ddresses
 - Map to physical track, cylinder, sector locations
 - Also, blocks are cached by the drive
 - | OS is unaware of physical locations
 - Reordering can reduce performance
 - Depending on placement and caching

$$LBA = ((C \times HPC) + H) \times SPT + S - 1$$

$$C = LBA \div (SPT \times HPC)$$

$$H = (LBA \div SPT) \bmod HPC$$

$$S = (LBA \bmod SPT) + 1$$

unaware /'ʌbə'weə/: không hay biết

Example: Disk Management

```
tiennd@tiennd:~$ sudo fdisk -l

Disk /dev/sda: 500.1 GB, 500107862016 bytes
255 heads, 63 sectors/track, 60801 cylinders, total 976773168 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk identifier: 0x404ccd9b

   Device Boot      Start         End      Blocks   Id  System
/dev/sda1            356530606     976773119     310121257    f   W95 Ext'd (LBA)
/dev/sda2            * 146801970     356530544     104864287+    7   HPFS/NTFS/exFAT
/dev/sda3                2048     130070527      65034240   83   Linux
/dev/sda4            130070528     146800639      8365056    82   Linux swap / Solaris
/dev/sda5            356530608     482351624      62910508+    7   HPFS/NTFS/exFAT
/dev/sda6            482367488     976773119     247202816   83   Linux
```

- ❑ Disk size = <sector num>*<sector size>
= 976773168 * 512 = 500107862016=465GB

Pitfall: Peak Performance

- ❑ Peak I/O rates are nearly impossible to achieve
 - | Usually, some other system component limits performance
 - | E.g., transfers to memory over a bus
 - Collision with DRAM refresh
 - Arbitration contention with other bus masters
 - | E.g., PCI bus: peak bandwidth ~133 MB/sec
 - In practice, max 80MB/sec sustainable



6.12. Concluding Remarks

- ❑ I/O performance measures
 - | Throughput, response time
 - | Dependability and cost also important
- ❑ Buses used to connect CPU, memory, I/O controllers
 - | Polling, interrupts, DMA
- ❑ I/O benchmarks
 - | TPC, SPECSFS, SPECWeb
- ❑ RAID
 - | Improves performance and dependability