

# SYSTEM SOFTWARE

## SYSTEM CALL

*Nguyen Huu Duc*

School of Information and Communication Technology  
Hanoi University of Science and Technology

# Overview of System Call

- In a computer system, applications are normally not allowed to access system resources; instead, they tell the operating system to help them access necessary resources.
  - System resources are balanced among applications
  - Protect critical data in the system from unwanted changes made by applications
  - Guarantee that an application only operates inside permitted space
- The operating system provides a set of system calls to satisfy the demands of the application that needs to access resources
  - Cover all types of resources
  - Constructed in a stable and sophisticated manner

- Modern processors support many different execution modes, two of which are used by the operating system
  - User mode: used when executing code in an application
    - only allowed to access the application's address space
    - only permitted to execute a limited set of instructions
    - be forbidden from direct access to data structures that are under protection by the operating system
  - Kernel mode: used when executing the operating system's instructions (such as system calls)
    - Be allowed to execute certain special instructions
    - Be allowed to access protected address spaces
    - Be allowed to run system programs

# System call and library functions

- System calls are performed in kernel mode, but library functions run in user mode (possibly require system calls)
- System calls that are not linked to applications
- Be limited in number (about 250)
- Library functions are rarely activated with call/return, while system calls make use of software interrupt
- For the ease of use, system calls are packed in a wrapper, which will then be used by applications as normal library functions

# System calls (1)

- File and I/O operations
  - `open()`, `close()`: Open and close file
  - `creat()`, `unlink()`: Create and delete file
  - `read()`, `write()`: Read and write file
  - `lseek()`: Moving pointer on file
  - `chmod()`: Change file access permissions
  - `mkdir()`, `rmdir()`: Create and delete folders

# System calls (2)

- File and I/O operations

- `stat()`: Read information about file status
- `access()`: Check file access permissions
- `chmod()`: Change file access permissions
- `chown()`: Change file owner
- `ioctl()`: Control devices

# System calls (3)

- Process management

- `fork()`: Create a new process
- `exec()`: Run an executable file
  - `execl(fn, arg0[, arg1, ..., argn], NULL)`
  - `execle(fn, arg0[, arg1, ..., argn], NULL, env)`
  - `execv(fn, argv)`
  - `execve(fn, argv, env)`
- `exit()`: Terminate a process
- `wait()`: Wait for child processes to finish

# System calls (4)

- Inter-process communication

- `kill()`, `sigsend()`: Send a signal to a process
- `pause()`: Suspend a process until receiving a required signal
- `sigaction()`: Set up procedure to handle signals



# System calls (5)

- Inter-process communication
  - `pipe()`: Create a pipeline
  - `socket()`: Create a socket
  - `bind()`: Assign an ID to a socket
  - `connect()`: Initiate a socket
  - `listen()`: Listen for data from a socket
  - `accept()`: Accept connection
  - `send()`: Send data to a socket
  - `recv()`: Receive data from a socket

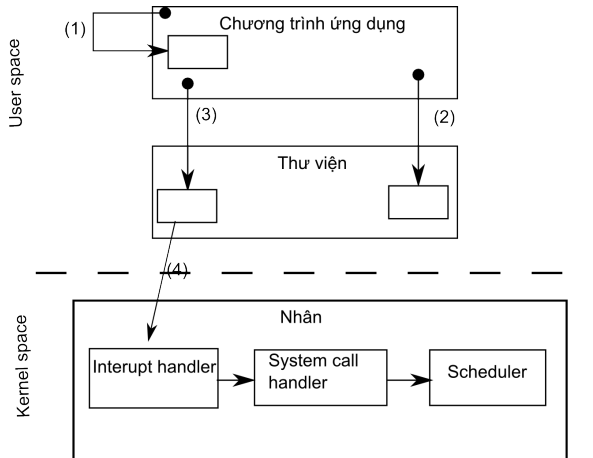
# System calls (6)

- Memory management
  - `brk()`, `sbrk()`: Change heap size

# Look up system calls

- `man 2`: look up a system call
- `man 3`: look up a library
- `man 4`: look up a device file
- `man 5`: look up file format
- `man 8`: look up command

# The process of making system calls



# The process of making system calls

- 1 Applications call routines inside their program
- 2 Application call functions in the library (without making system calls)
  - Example: `sscanf()`, `memset()`, `tolower()`,...
- 3 System programs call functions in the library (using system calls)
  - Example: `getchar()`, `printf()`,...
- 4 Library functions activate system calls

# Activate system calls

- Library functions (user mode) activate system calls (kernel mode) through interrupts
  - 0x21: MS-DOS
  - 0x80: Linux
  - 0x81: XEN
- A system call is determined with a number and is passed through register `eax`
- Parameters required by the system call are passed through general registers
  - In case the number of parameter is so large, these parameters are passed through a pointer
  - Pay attention to the parameters
- Variable `errno` stores the error code returned by the system call

# Activate system calls

- The interrupt handlers (`system_call`) choose the system call from a table of system calls (`sys_call_table`)
  - Store general registers (also parameters of the system call) on kernel stack
  - Choose system call from `sys_call_table` and make a request for the corresponding handler
  - Return control to the library function through instruction `iret`

```
system_call:
    SAVE_ALL
    ...
    call *sys_call_table(0, %eax, 4)
    ...
ret_from_sys_call:
    ...
    iret
```

- Some system calls cannot terminate immediately
  - nanosleep, read, write
- The process need to be scheduled in order to be able to pause (and transfer system control to another process) and wait for success signals from hardware




# Some system calls

```
int setuid (uid_t uid)
```

- Set valid UID for the current process
- Used frequently in server applications such as `httpd`, `bind`,...
- `uid`: user identifier
- Return 0 if succeed, -1 if fail(`errno=ERPERM`)

# setuid



## 1: An application makes system calls

- Call conventions `cdecl`
- Instruction usage `call`
- Return address on stack
- Transfer control to library functions

## 2: A library function prepares to make system call

- Check parameters
- Trigger 0x80 interrupt
  - Switch to kernel mode
  - SS,ESP, EFLAGS, CS, EIP is pushed on the kernel stack
  - Assign values SS,ESP, EFLAGS, CS, EIP to kernel

## 3: Prepare to make system calls

- The identifying number of the system call
- Store values that are currently on general register (es, ds, eax, ebp, edi, esi, edx, ecx, ebx)
- Pass es, ds to kernel data space

<- Ret addr

```
24: * 0 (% esp) -% ebx first argument
25: * 4 (% esp) -% ecx second argument
26: * 8 (% esp) -% edx third argument
27: * C (% esp) -% esi fourth parameter
28: * 10 (% esp) -% edi fifth argument
29: * 14 (% esp) -% ebp
30: * 18 (% esp) -% eax
31: * 1C (% esp) -% ds
32: * 20 (% esp) -% es
33: * 24 (% esp) -% user_eax
34: * 28 (% esp) -% user_eip
35: * 2C (% esp) -% user_cs
36: * 30 (% esp) -% user_eflags
37: * 34 (% esp) -% user_esp
38: * 38 (% esp) -% user_ss
```

## 3: Prepare to make system calls

- Check the identifying number of the system calls
- Call `sys_setuid`
  - `call *sys_call_table (0,%eax, 4)`

## 4: Execute a system call

- Check parameters
  - EAGAIN: Resources are already in use
  - EPERM: Permission denied
- Execute effective uid



## 5: Prepare to return results to the library function

- Put the value of `eax` to the 24th position (`%esp`)
- Restore the previous values of general registers
- Restore the value of `eax`
- `iret`
  - Restore `eip`, `cs`, `eflags`, `esp`, `ss`
  - Switch to user mode

## 6: Prepare to return results to the application

- Check for errors
- Set up errno

## 7: The application

- Receive returned results through eax
- Check for error (errno)

# Steps in constructing a system call

- Prepare tools for re-compiling the kernel of the operating system
- Prepare source code of the kernel of the operating system
- Change and add code of the new system call
- Recompile the kernel of the operating system
- Test the new system call

# Construct a system call (1)

- Source code of the system call
  - /usr/src/linux/mycall.c

mycall.c

```
#include<linux/linkage.h>
asm linkage long sys_mycall(int i)
{
    return i+10;
}
```

## Construct a system call (2)

- Add the system call to `syscall_table.S`
  - `/usr/src/linux/arch/i386/kernel/syscall_table.S`

`mycall.c`

```
...  
.long sys_mycall
```

## Construct a system call (3)

- Add the system call to `unistd.h`
  - `/usr/src/linux/include/asm-i386/unistd.h`

`unistd.h`

```
...  
#define __NR_vmsplice 316  
#define __NR_mycall 317  
...  
#define NR_syscalls 318
```

# Construct a system call (4)

- Add the system call to syscalls.h
  - /usr/src/linux/include/linux/syscalls.h

syscalls.h

...

```
asmlinkage long sys_mycall(int i);
```

# Construct a system call (5)

- Change Makefile of the kernel
  - Create Makefile for the "mycall" folder
  - Change Makefile of the kernel
    - Insert "mycall/" to "core-y"

## Makefile

```
obj-y := mycall.o
```